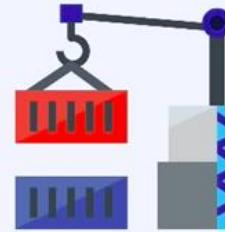




**kubernetes**

# Official definition of Kubernetes



- ▶ Open source container **orchestration tool**
- ▶ Developed by **Google**
- ▶ Helps manage containerized applications  
in **different deployment environments**

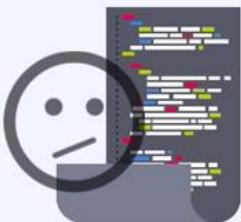


**What problems does Kubernetes solve?**

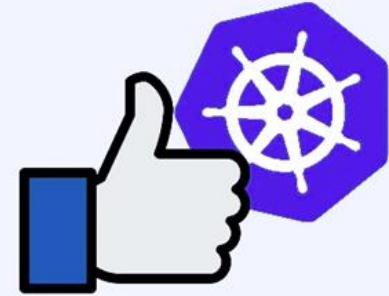
**What are the tasks of an orchestration tool?**

# Need for container orchestration tool

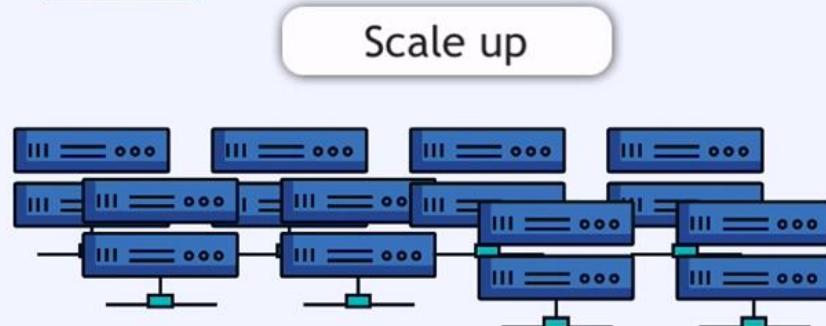
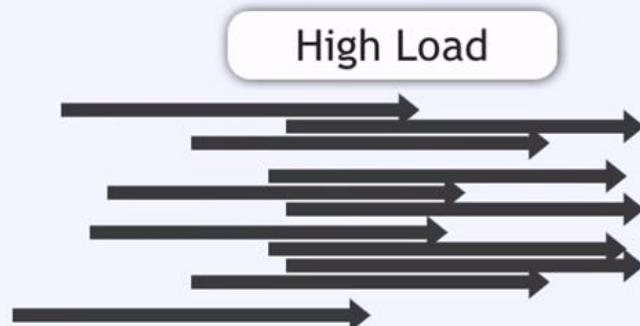
- ▶ Trend from **Monolith** to **Microservices**
- ▶ Increased usage of containers
- ▶ Demand for a **proper way of managing** those hundreds of containers



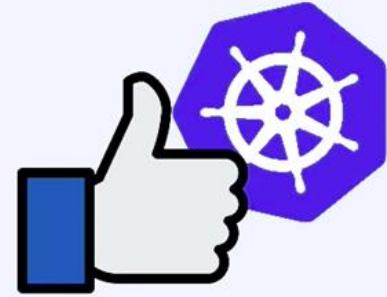
# What features do orchestration tools offer?



- ▶ High **Availability** or no downtime
- ▶ **Scalability** or high performance



# What features do orchestration tools offer?



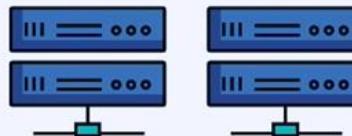
- ▶ High **Availability** or no downtime
- ▶ **Scalability** or high performance



Load decreasing..

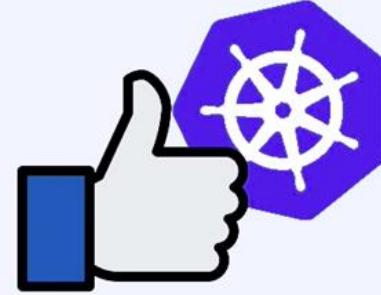


Scale down



# What features do orchestration tools offer?

- ▶ High **Availability** or no downtime
- ▶ **Scalability** or high performance
- ▶ **Disaster recovery** - backup and restore



# Kubernetes

## Architecture



**Node** = virtual or physical machine



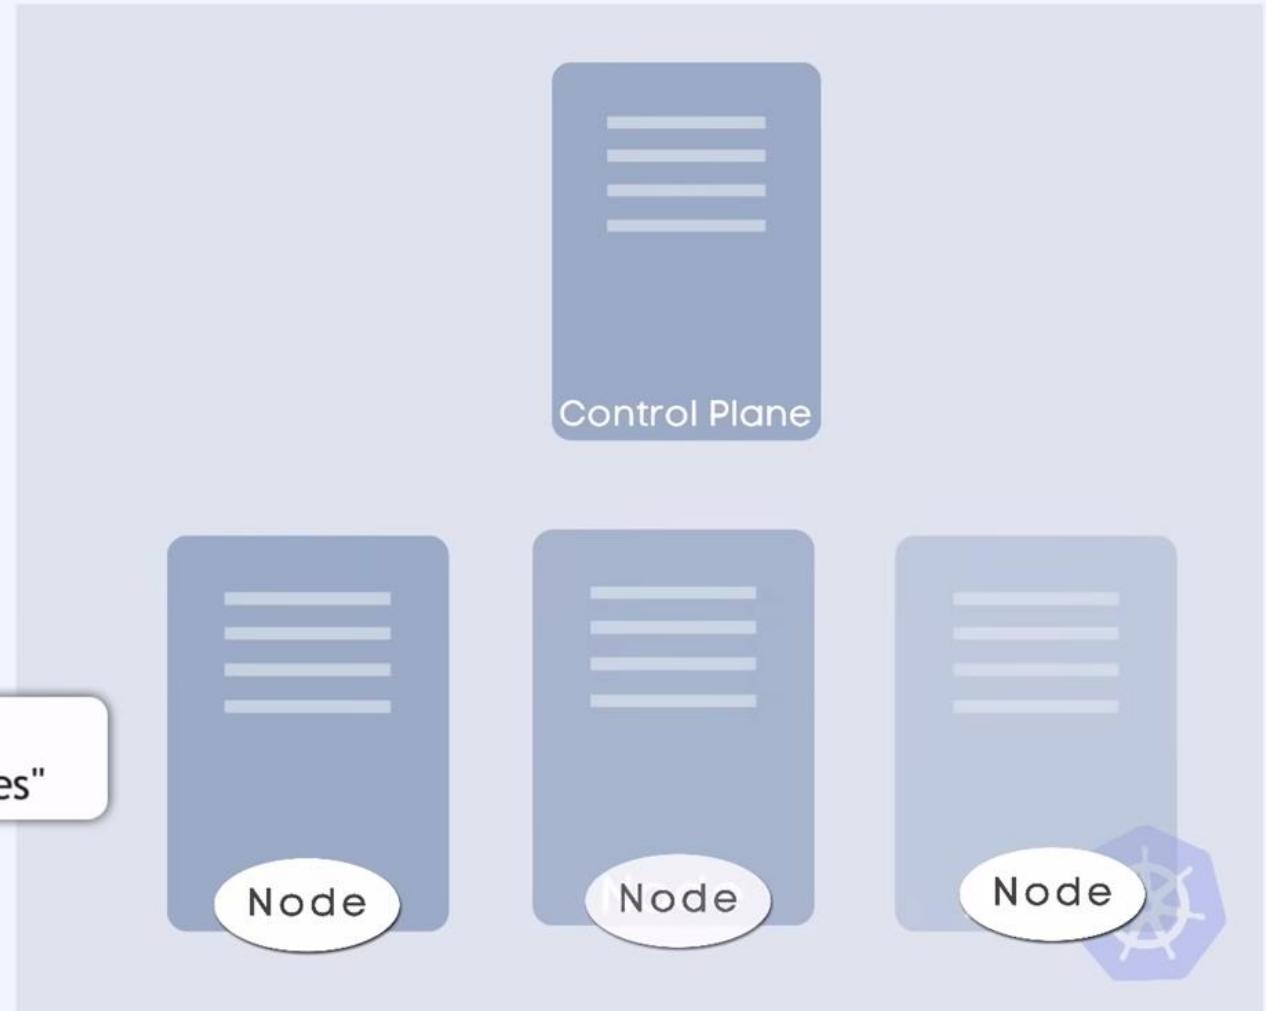
Control Plane



# Kubernetes

## Architecture

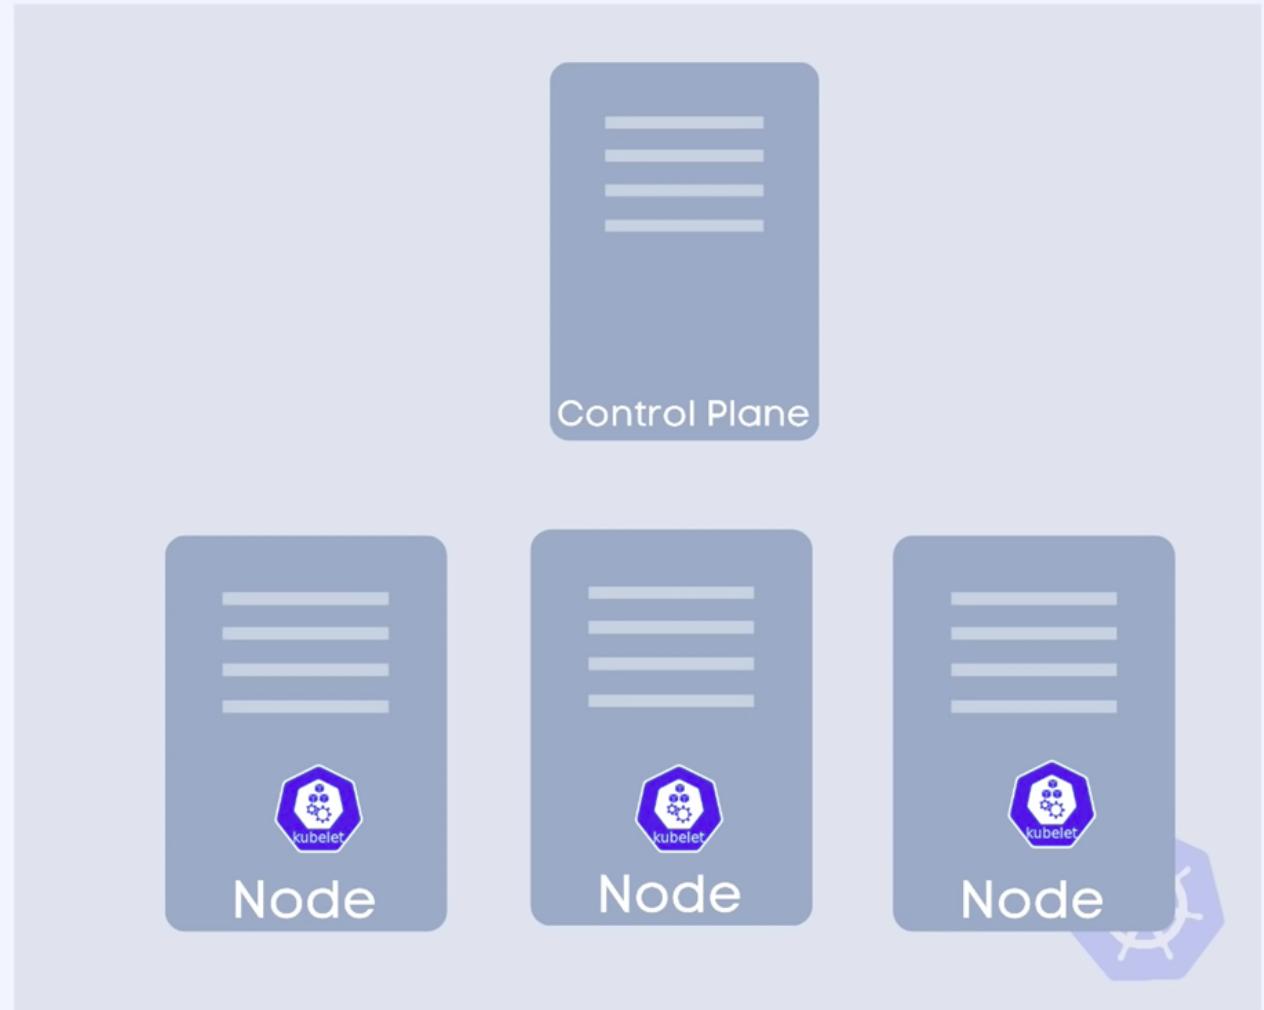
**Worker Nodes =**  
mostly just referred to as "Nodes"



# Kubernetes Architecture

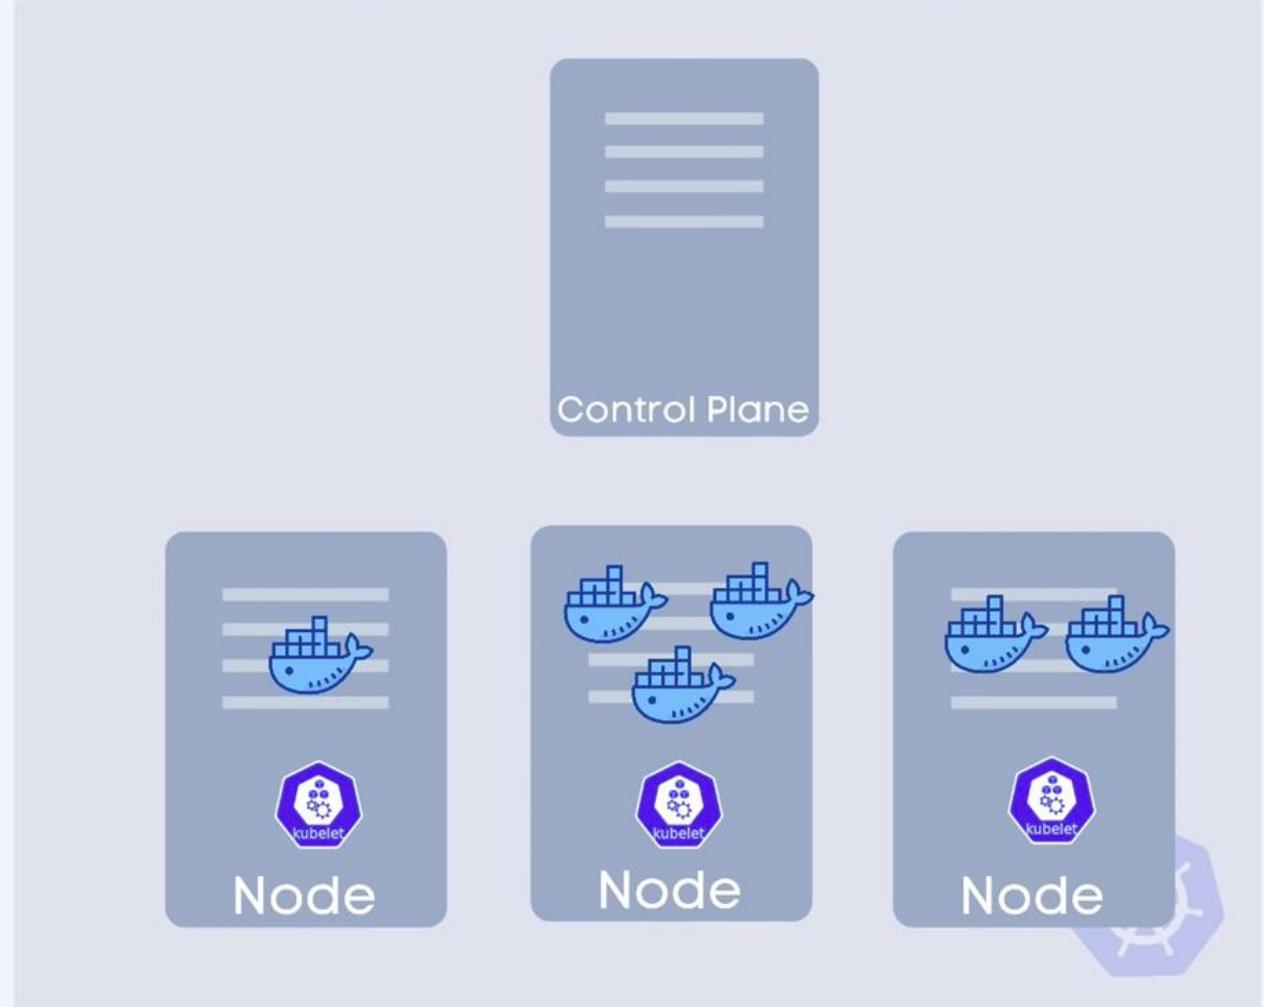


primary "node agent"



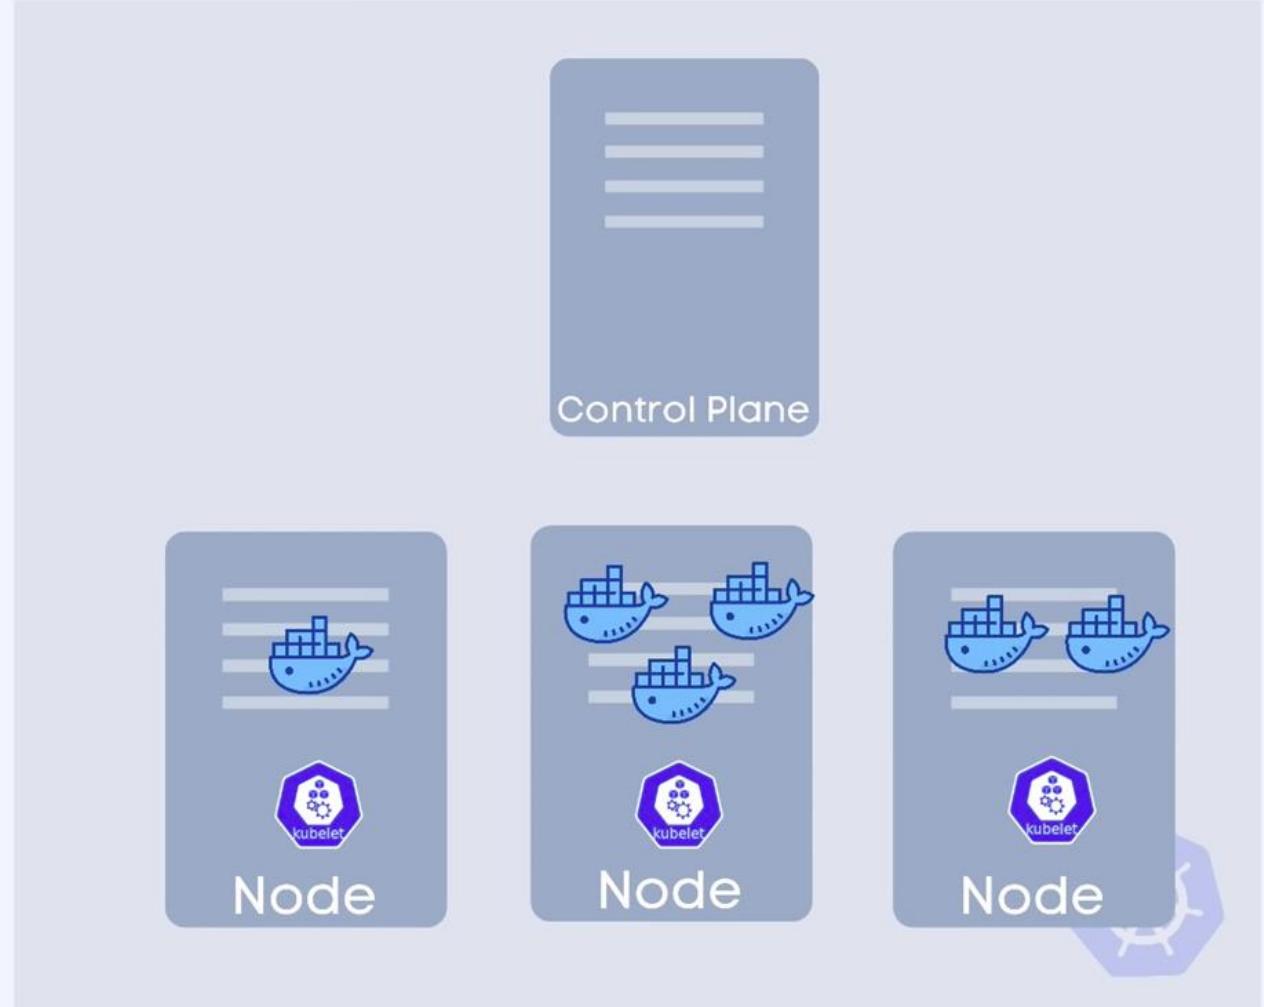
# Kubernetes

## Architecture



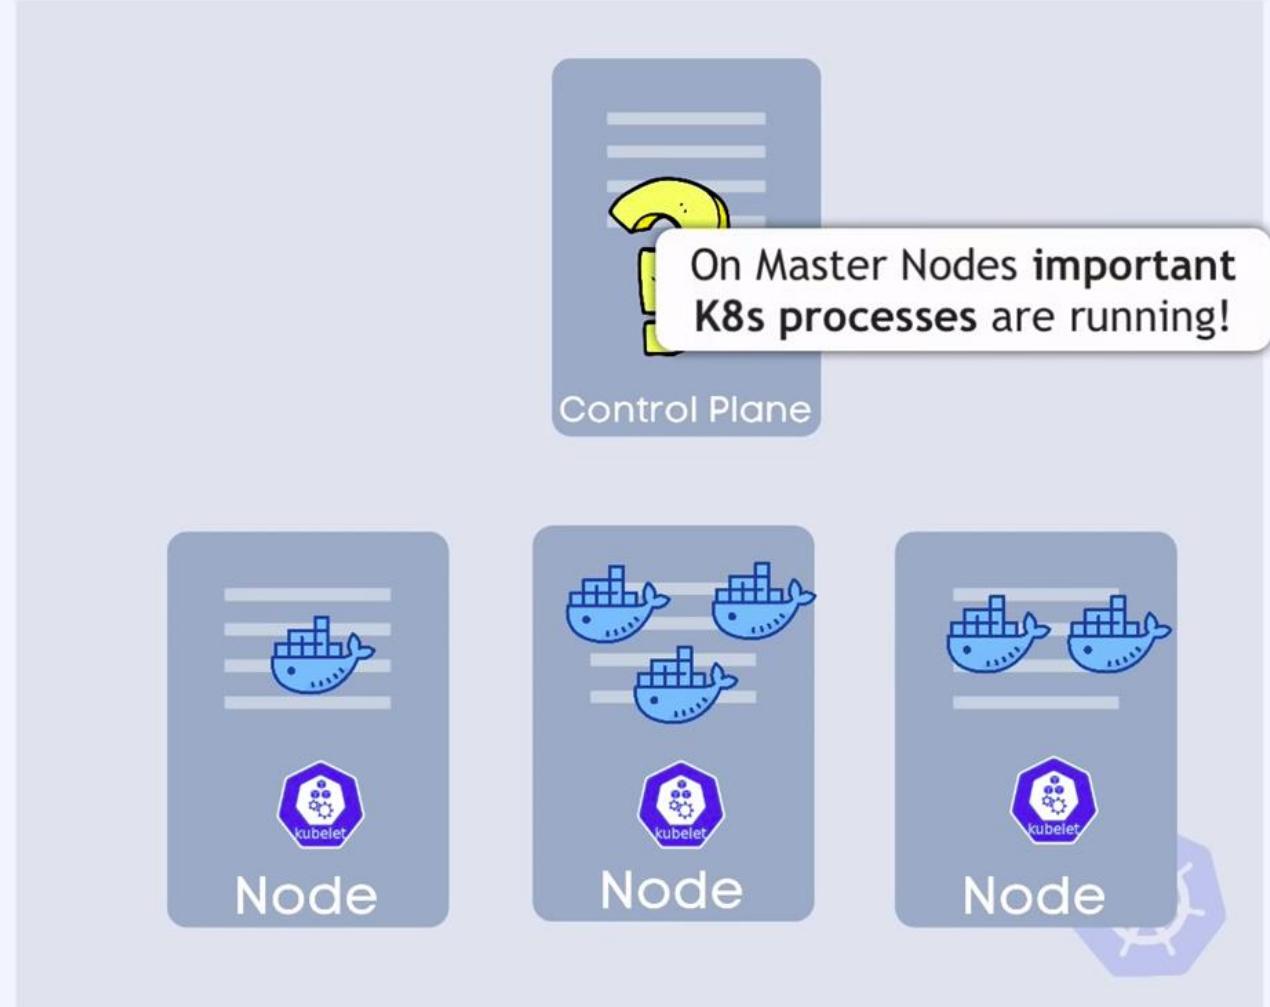
# Kubernetes Architecture

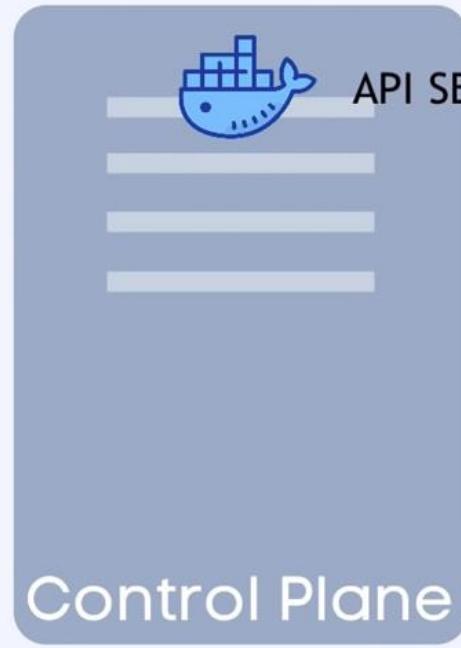
On Worker Nodes your applications are running



# Kubernetes

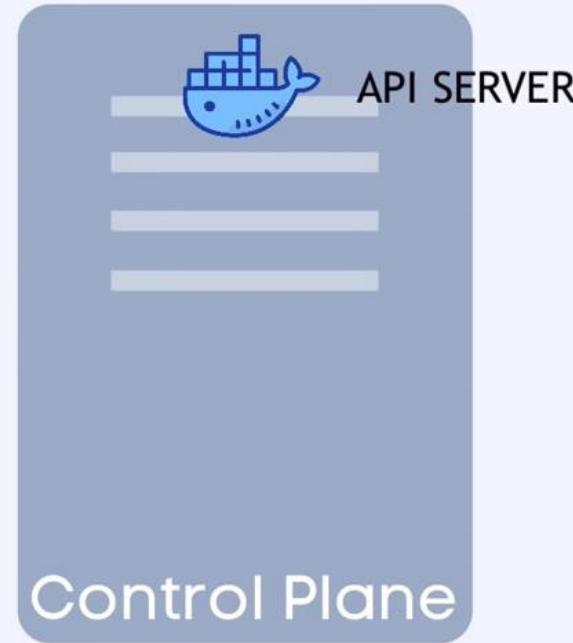
## Architecture







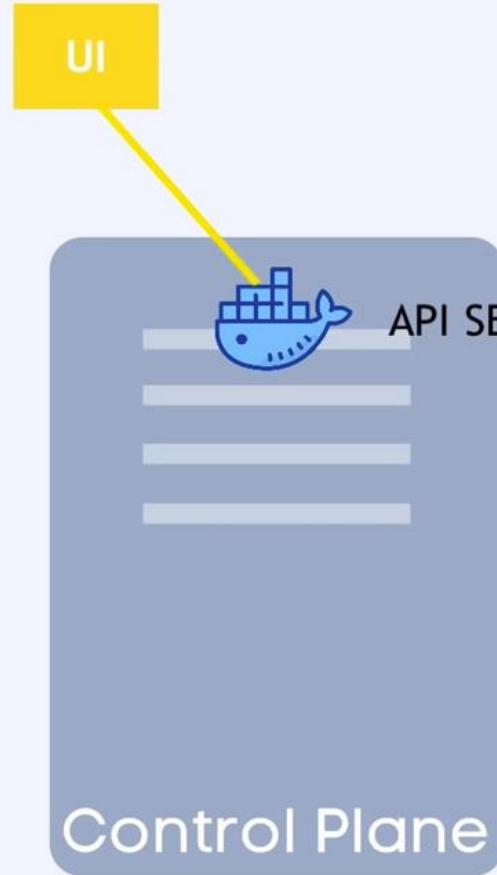
**API Server =**  
Entrypoint to K8s cluster





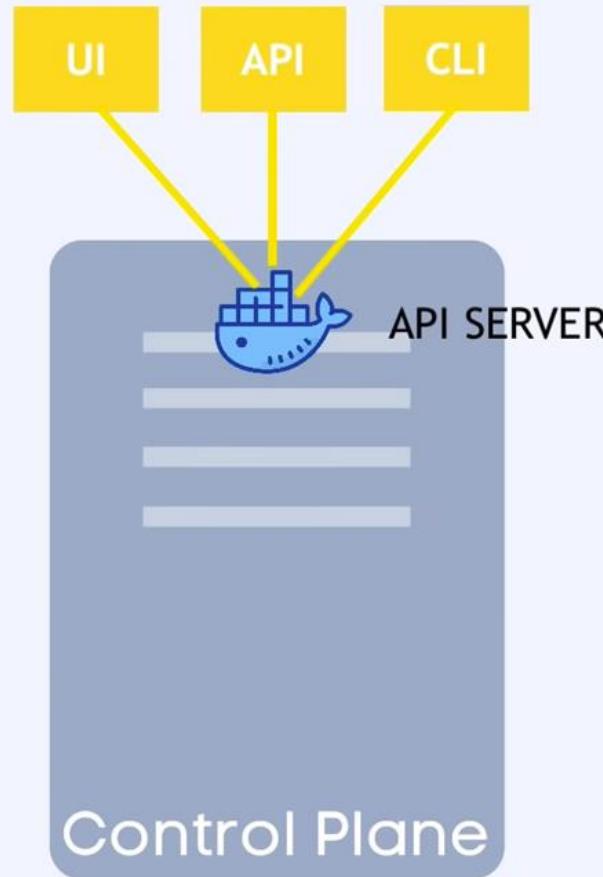
API Server =  
Entrypoint to K8s cluster

The screenshot shows the Kubernetes dashboard interface. On the left, a sidebar lists various cluster components: Cluster, Namespaces, Nodes, Persistent Volumes, Roles, Storage Classes, and Namespace (set to default). Under Workloads, it lists Cron Jobs, Daemon Sets, Deployments, Jobs, Pods, Replica Sets, Replication Controllers, and Stateful Sets. The main area displays "Workloads Statuses" with three green circular charts: Deployments (100.0%), Pods (100.0%), and Replica Sets (100.0%). Below this, the "Deployments" section shows a single entry for "nginx" with 1 pod, created 55 seconds ago, using the "nginx:1.15.5" image. The "Pods" section shows the same "nginx" pod with its status as "Running" on node "ak8-nodepool1-79590246-0".





**API Server** =  
Entrypoint to K8s cluster

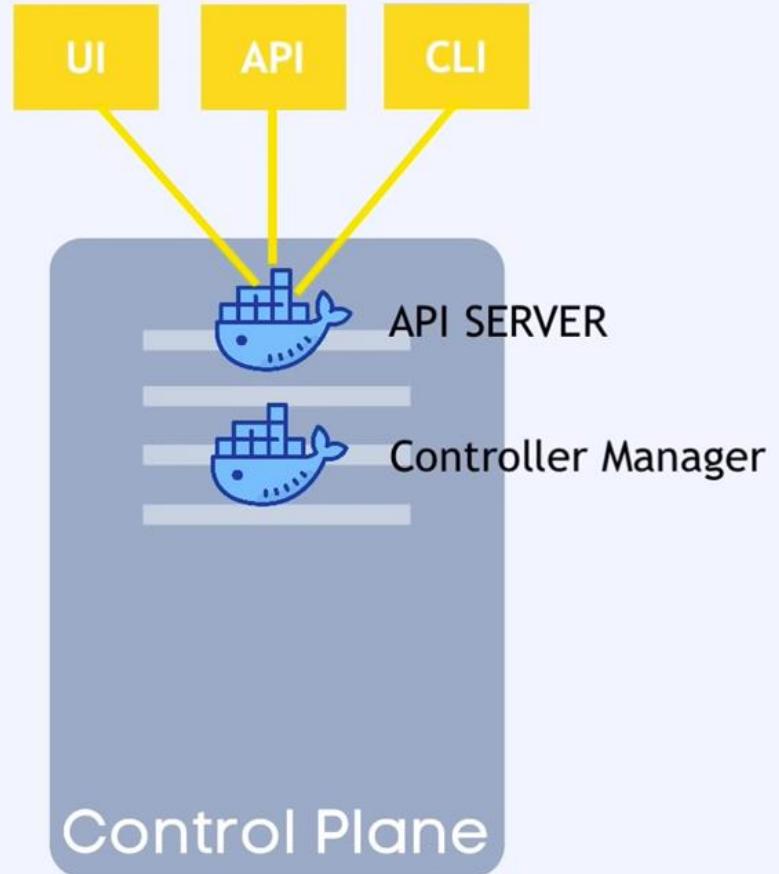




**API Server** =  
Entrypoint to K8s cluster



**Controller Manager** = keeps track of  
what's happening in the cluster





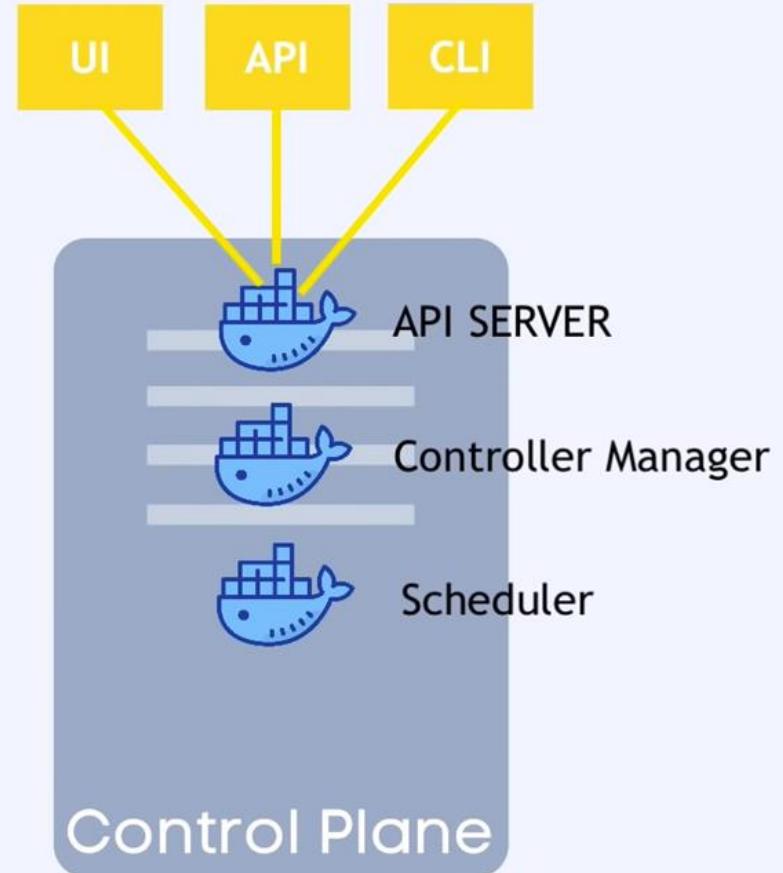
**API Server** =  
Entrypoint to K8s cluster



**Controller Manager** = keeps track of  
what's happening in the cluster



**Scheduler** =  
ensures Pods placement





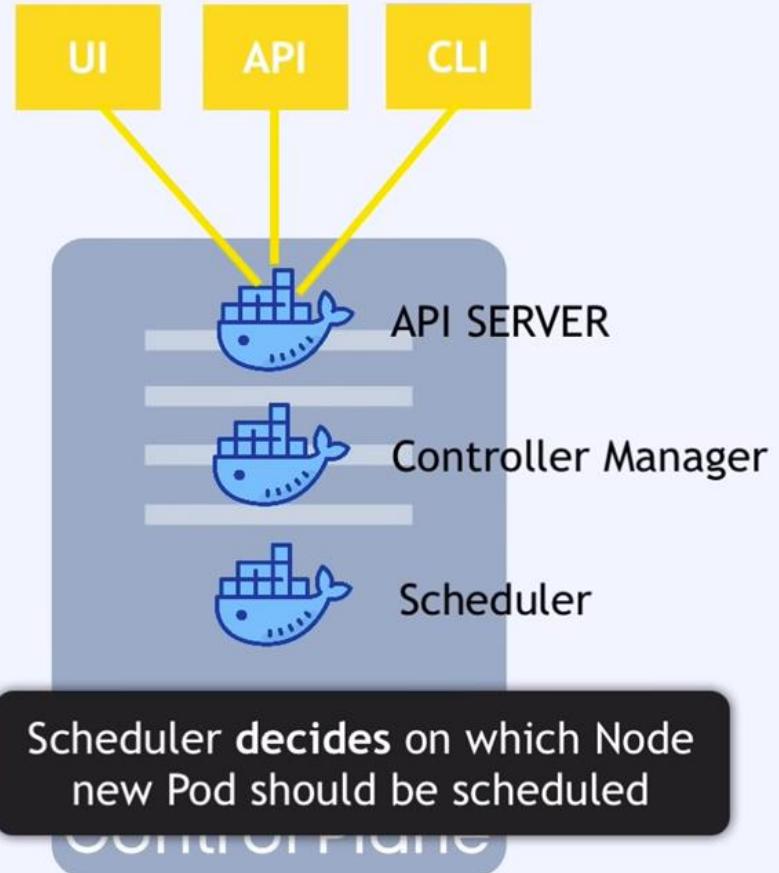
**API Server** =  
Entrypoint to K8s cluster



**Controller Manager** = keeps track of  
what's happening in the cluster



**Scheduler** =  
ensures Pods placement





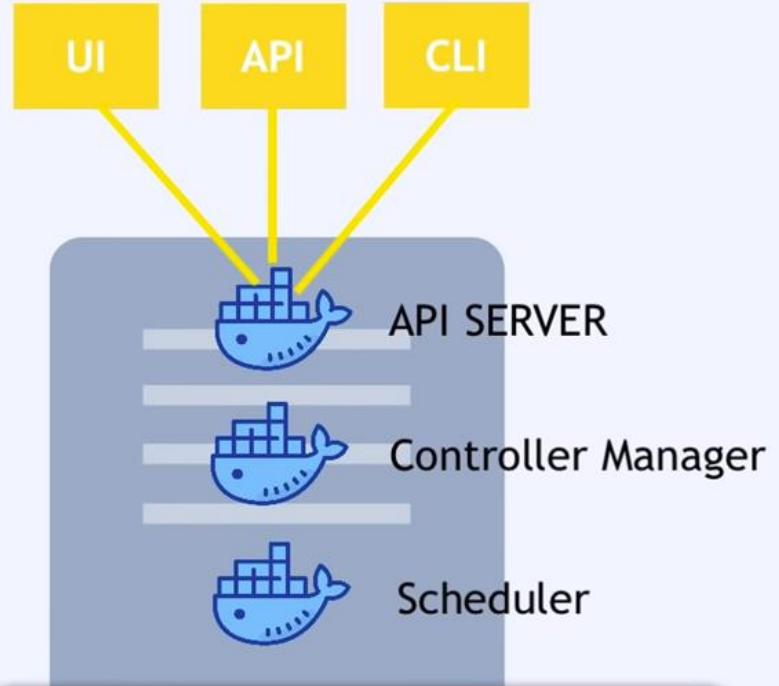
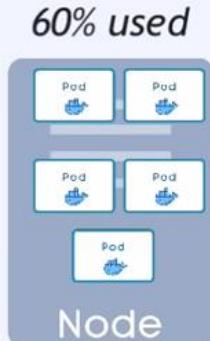
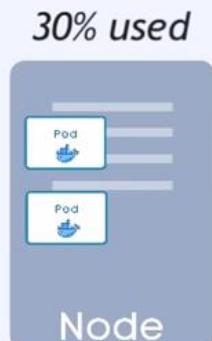
**API Server** =  
Entrypoint to K8s cluster



**Controller Manager** = keeps track of  
what's happening in the cluster



**Scheduler** =  
ensures Pods placement



Scheduler **decides** on which Node  
new Pod should be scheduled



**API Server** =  
Entrypoint to K8s cluster

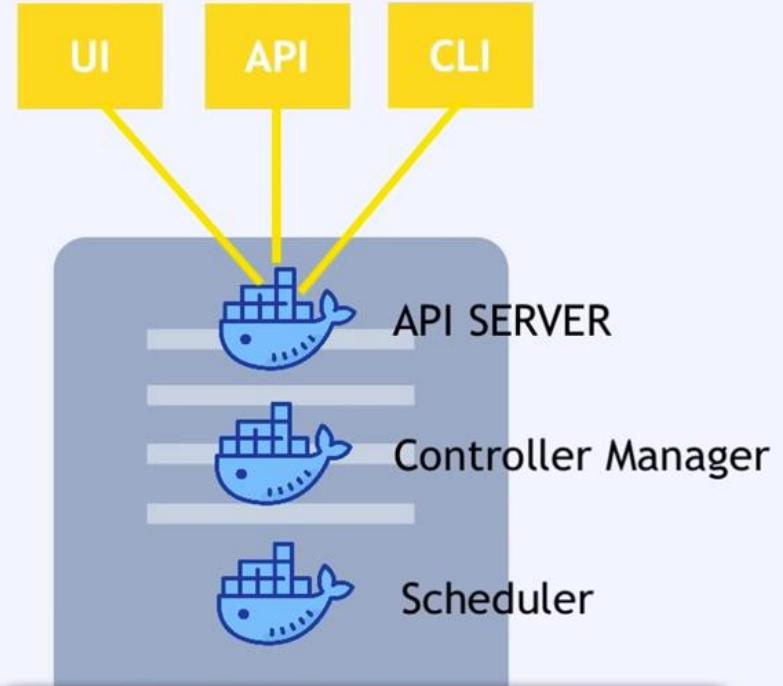
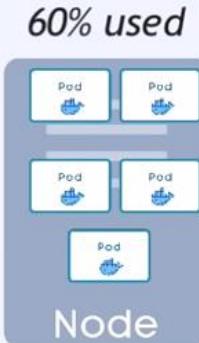


**Controller Manager** = keeps track of  
what's happening in the cluster



**Scheduler** =  
ensures Pods placement

Where to put the Pod?



Scheduler **decides** on which Node  
new Pod should be scheduled



**API Server** =  
Entrypoint to K8s cluster



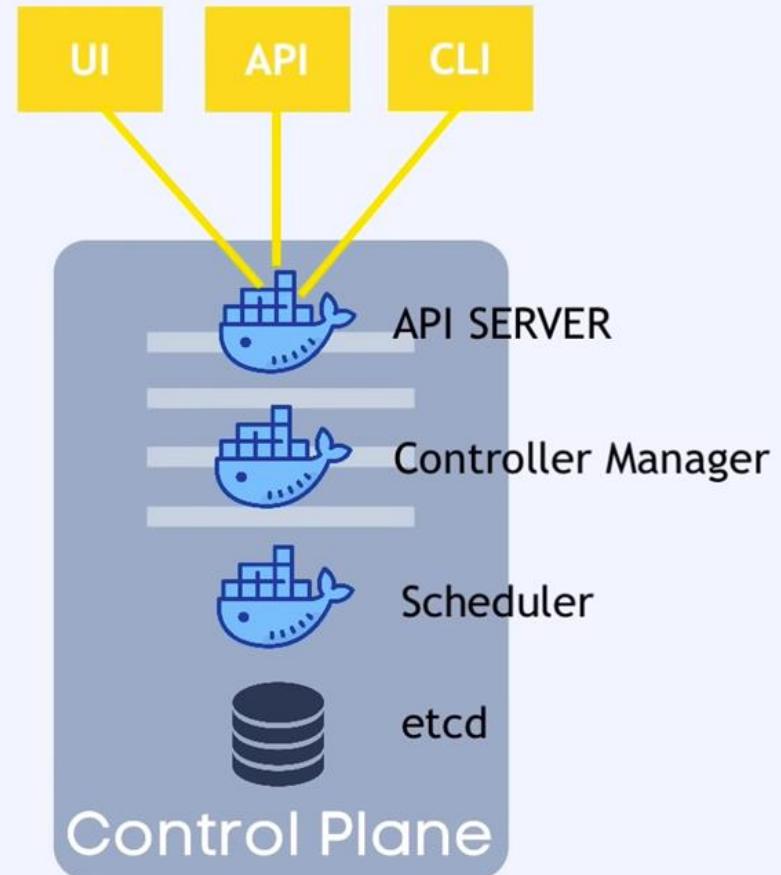
**Controller Manager** = keeps track of  
what's happening in the cluster

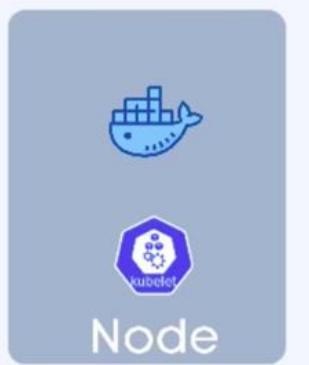
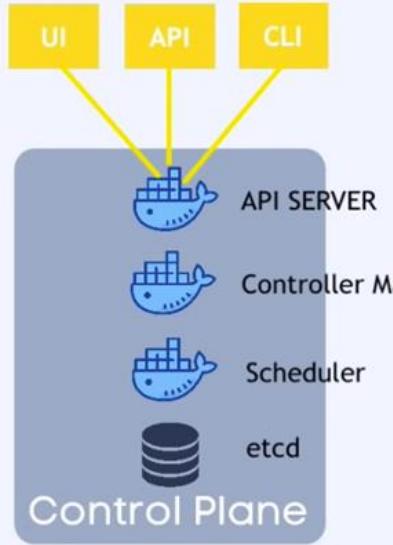


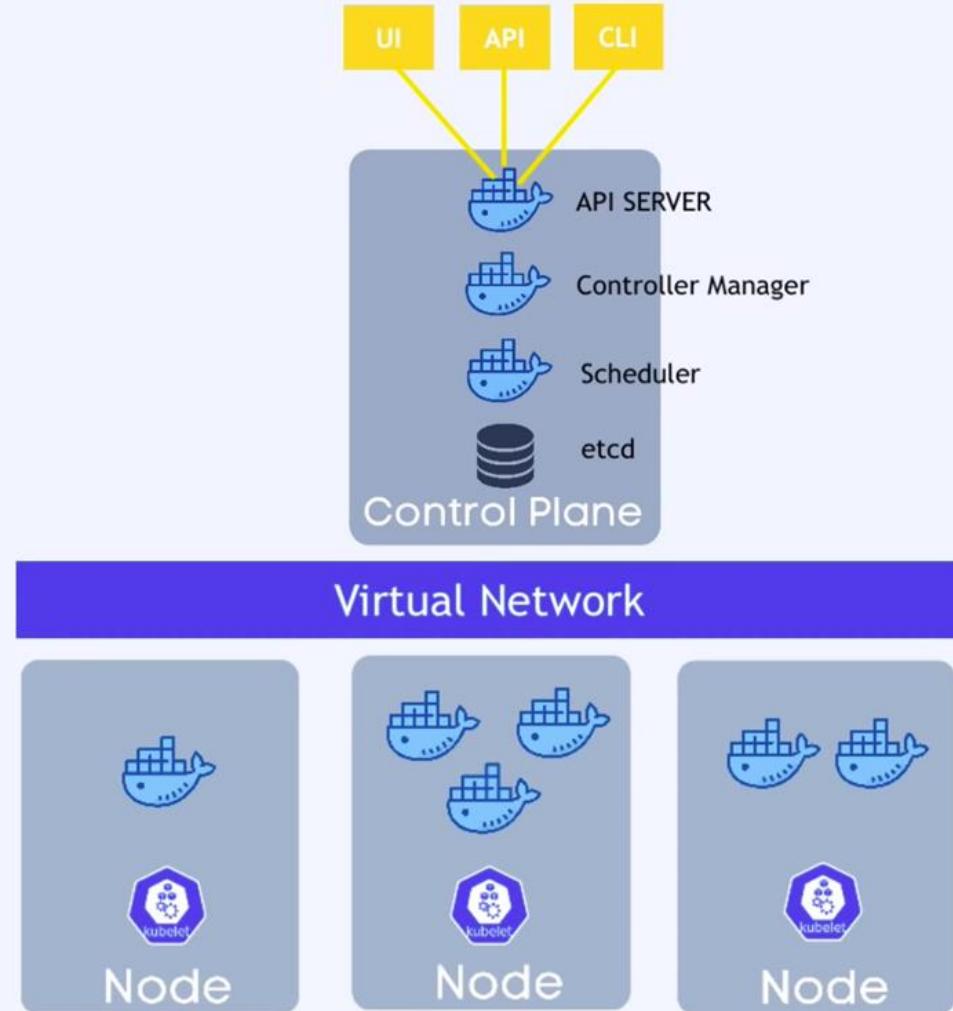
**Scheduler** =  
ensures Pods placement



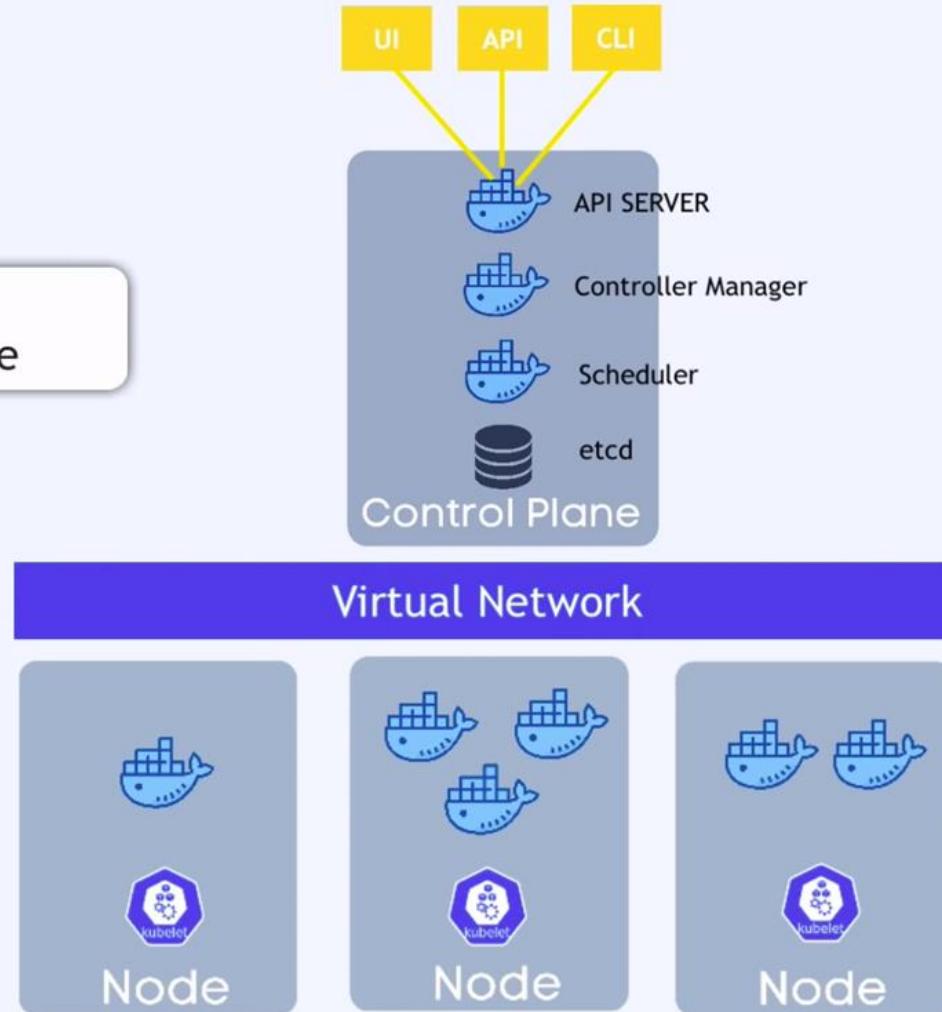
**etcd** =  
Kubernetes backing store

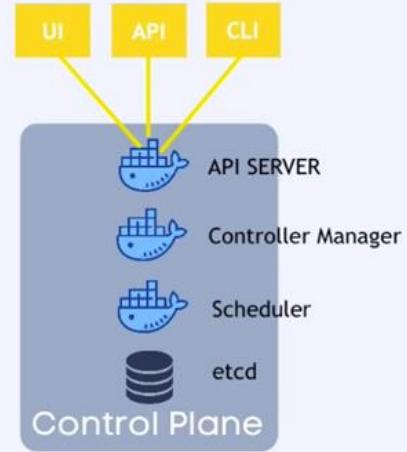






**Virtual Network =**  
Creates one unified machine





**Worker Nodes**

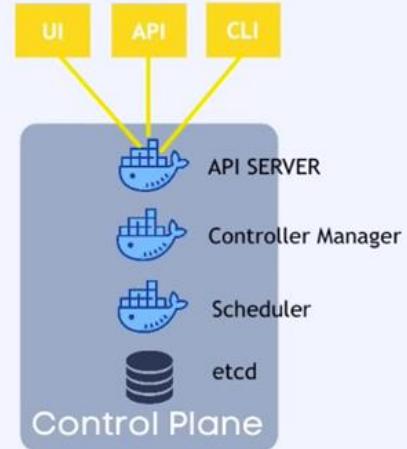
higher workload

much bigger and more resources

## Control Plane Nodes

handful of master processes

much more important



## Worker Nodes

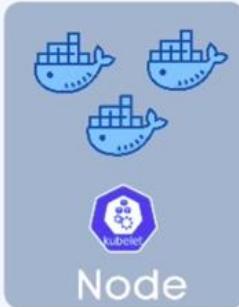
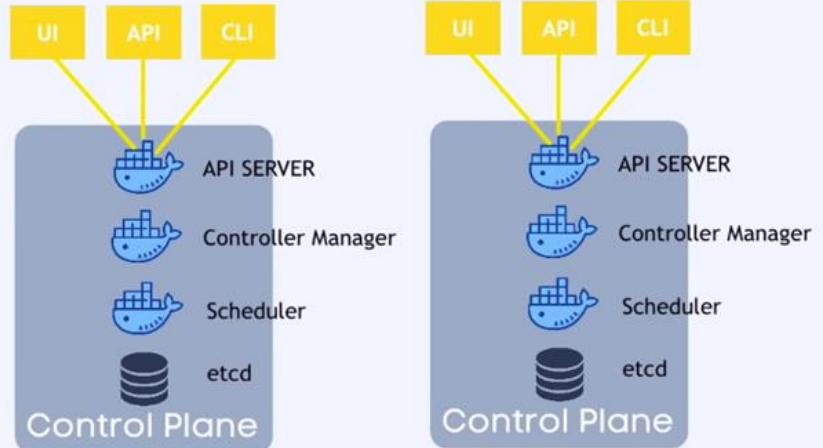
higher workload

much bigger and more resources

## Control Plane Nodes

handful of master processes

much more important



## Worker Nodes

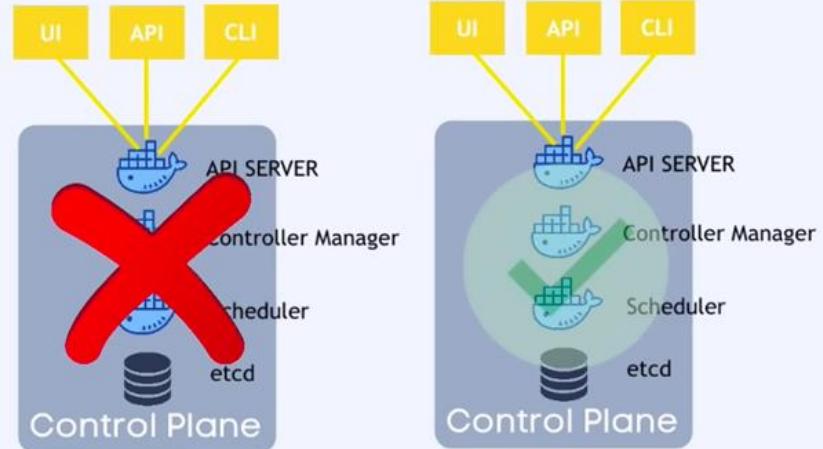
higher workload

much bigger and more resources

## Control Plane Nodes

handful of master processes

much more important



## Worker Nodes

higher workload

much bigger and more resources

# Main Kubernetes Components



# Main Kubernetes Components



Role of each component

Pod

ConfigMap

Service

Secret

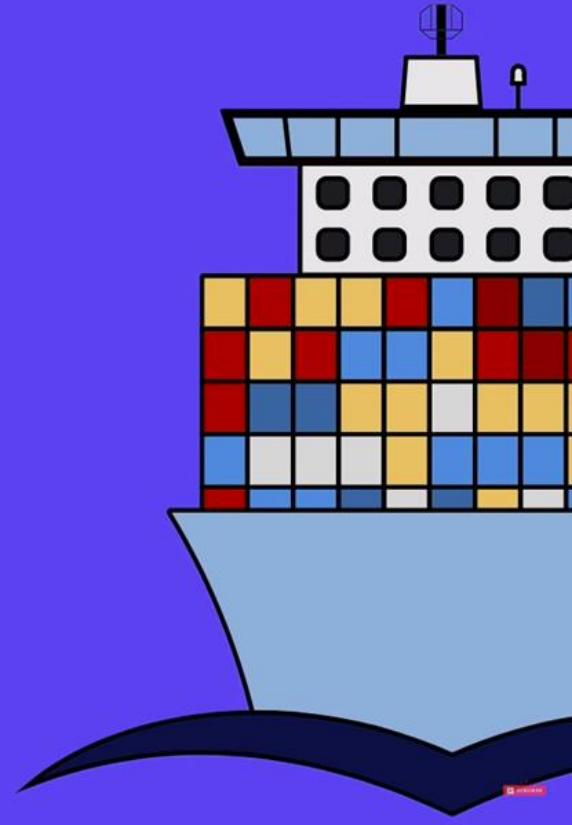
Ingress

Deployment

StatefulSet

DaemonSet

# Node and Pod



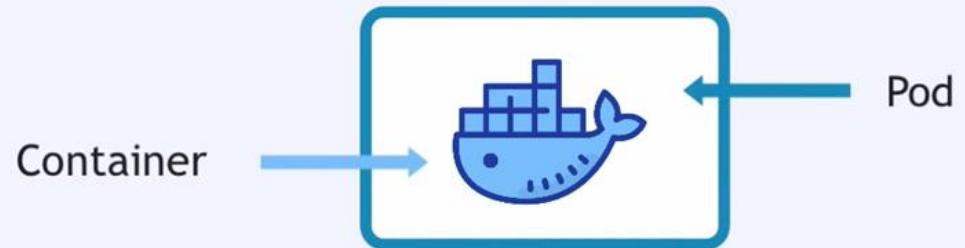
# Node

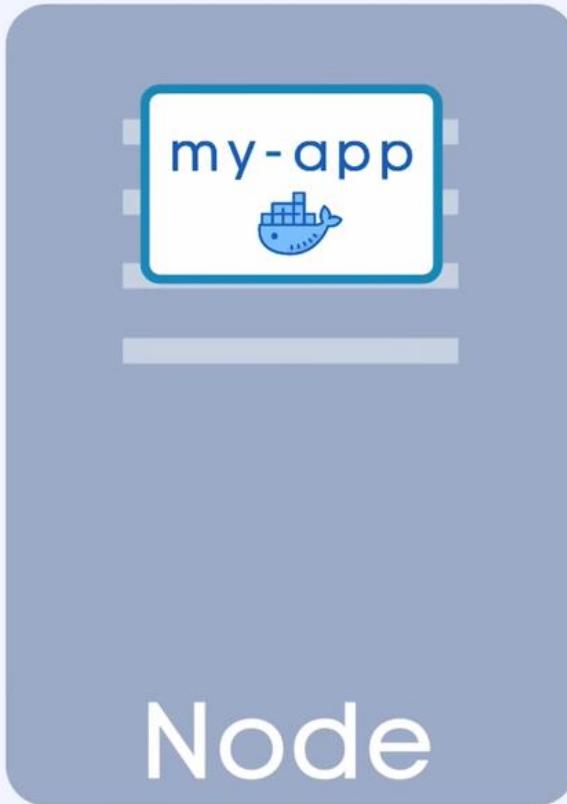


**Node** =  
virtual or physical machine

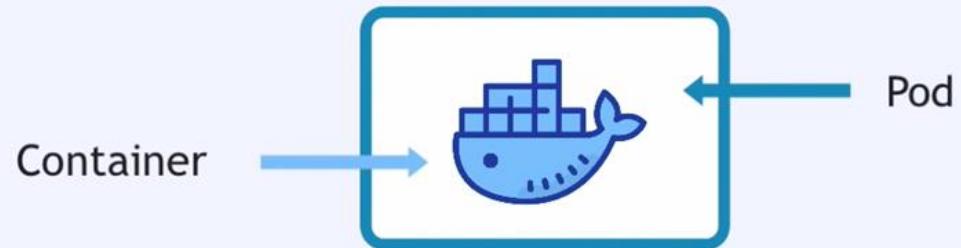


- ▶ **Smallest** unit in Kubernetes
- ▶ **Abstraction** over container

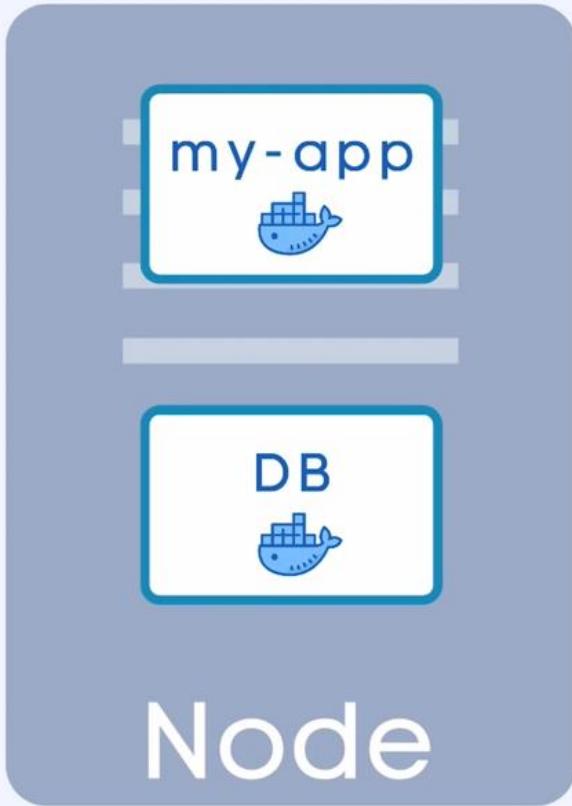




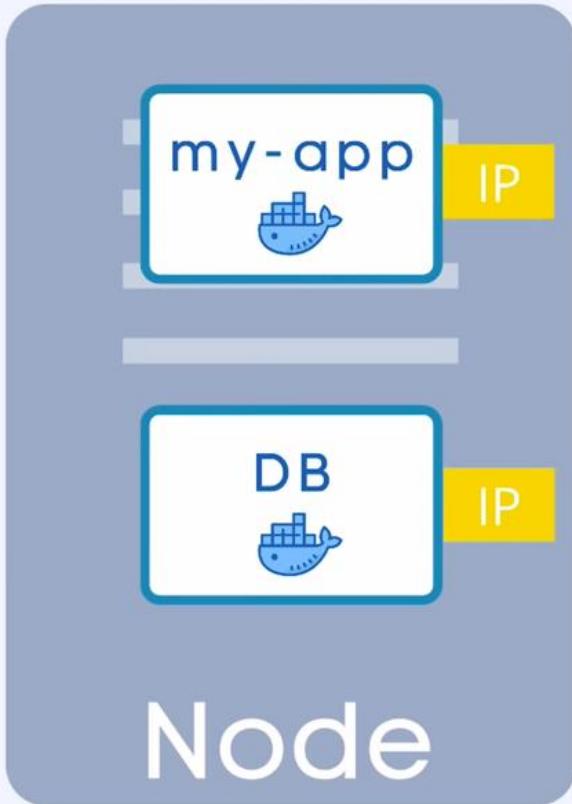
- ▶ **Smallest** unit in Kubernetes
- ▶ **Abstraction** over container



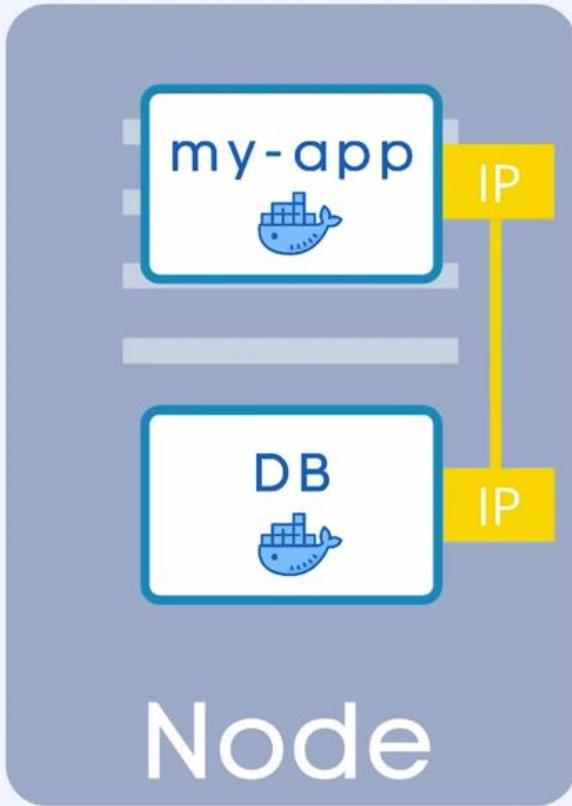
You only interact with the Kubernetes layer



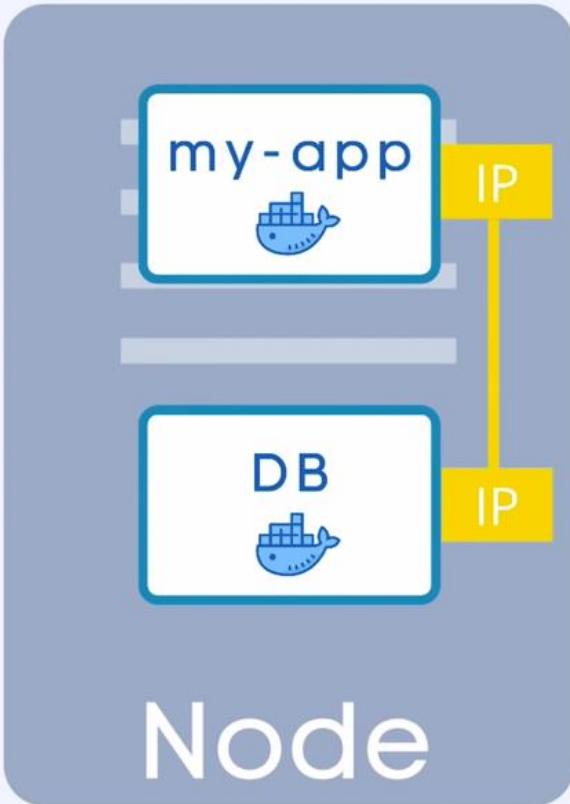
- ▶ **Smallest** unit in Kubernetes
- ▶ **Abstraction** over container
- ▶ Usually **1 Application** per Pod



- ▶ **Smallest** unit in Kubernetes
- ▶ **Abstraction** over container
- ▶ Usually **1 Application** per Pod



- ▶ **Smallest** unit in Kubernetes
- ▶ **Abstraction** over container
- ▶ Usually **1 Application** per Pod
- ▶ Each Pod gets its **own IP address**

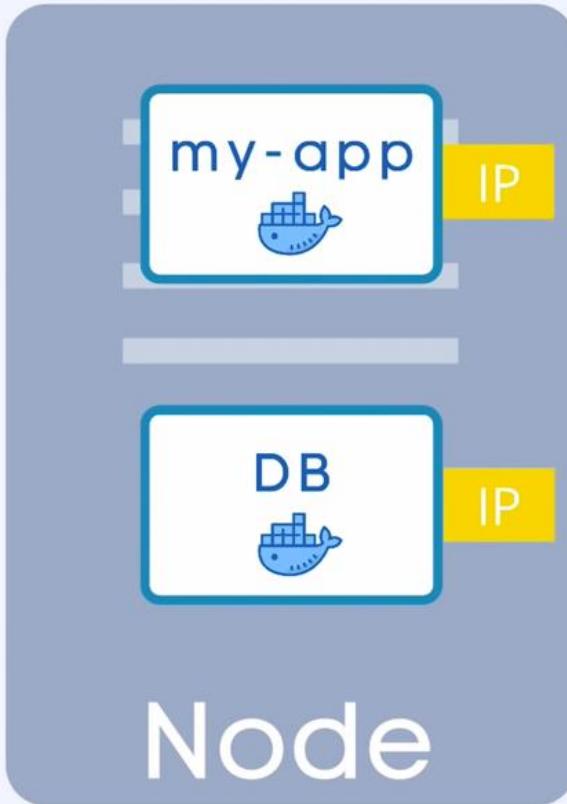


- ▶ **Smallest** unit in Kubernetes
- ▶ **Abstraction** over container
- ▶ Usually **1 Application** per Pod
- ▶ Each Pod gets its **own IP address**

Pods are **ephemeral!**



- ▶ **Smallest** unit in Kubernetes
- ▶ **Abstraction** over container
- ▶ Usually **1 Application** per Pod
- ▶ Each Pod gets its **own IP address**



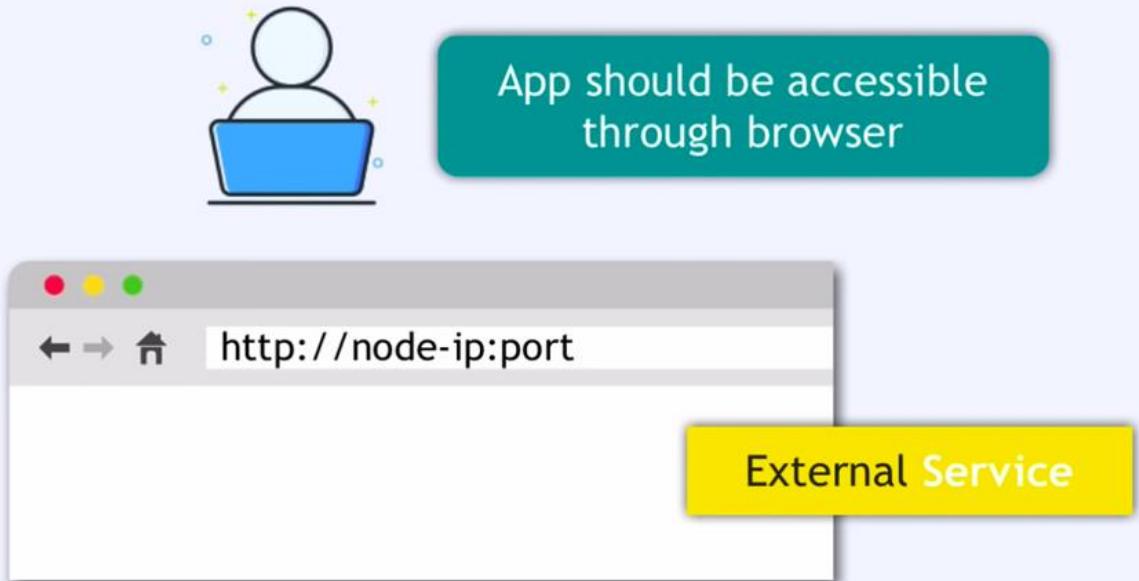
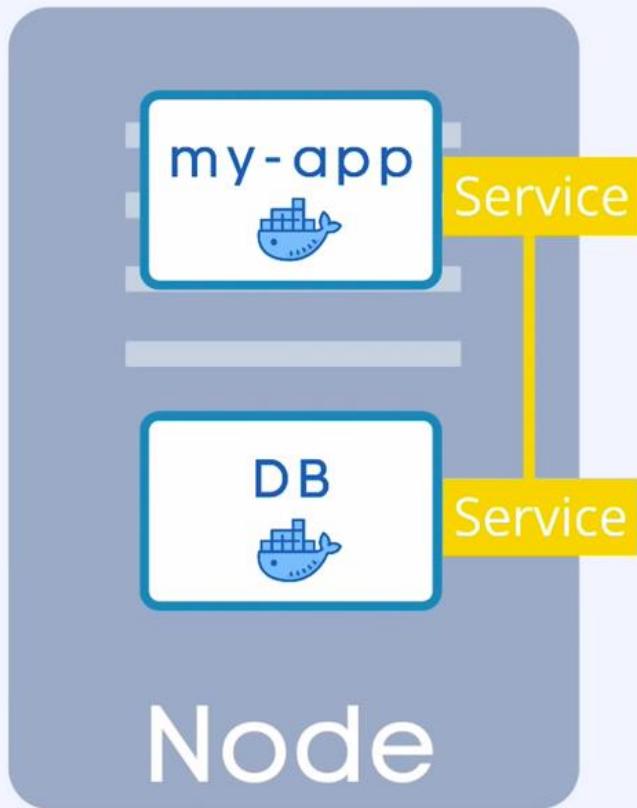
- ▶ **Smallest** unit in Kubernetes
- ▶ **Abstraction** over container
- ▶ Usually **1 Application** per Pod
- ▶ Each Pod gets its **own IP address**
- ▶ **New IP address** on re-creation

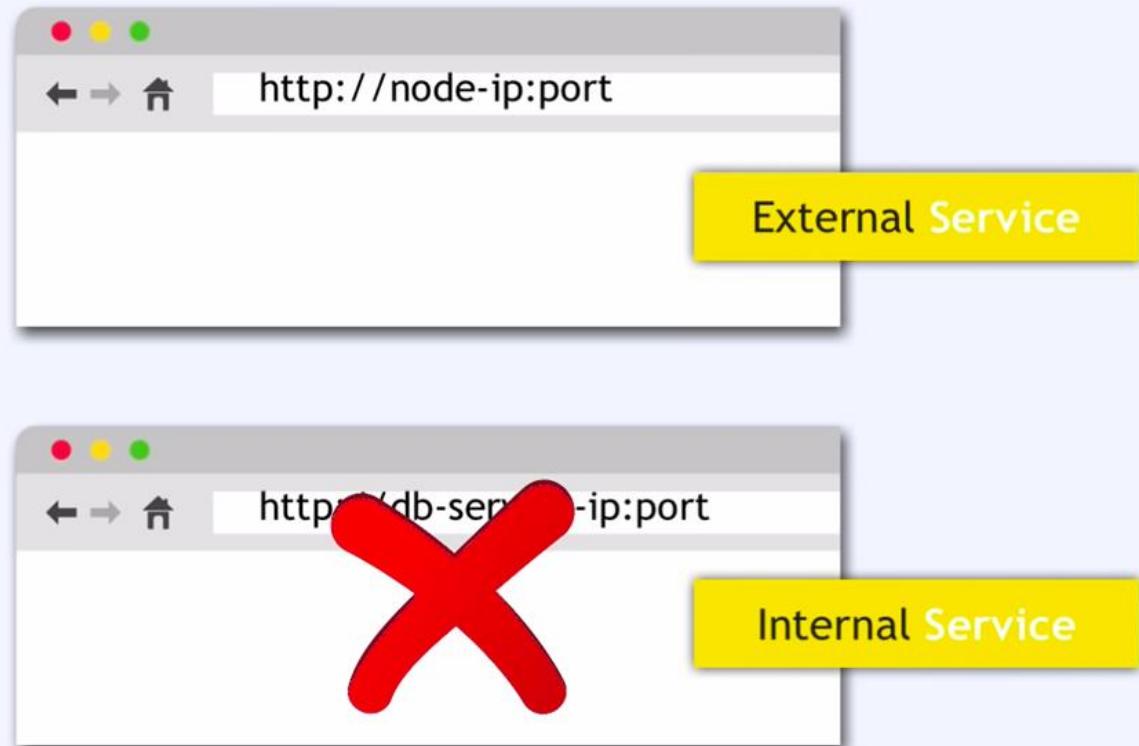
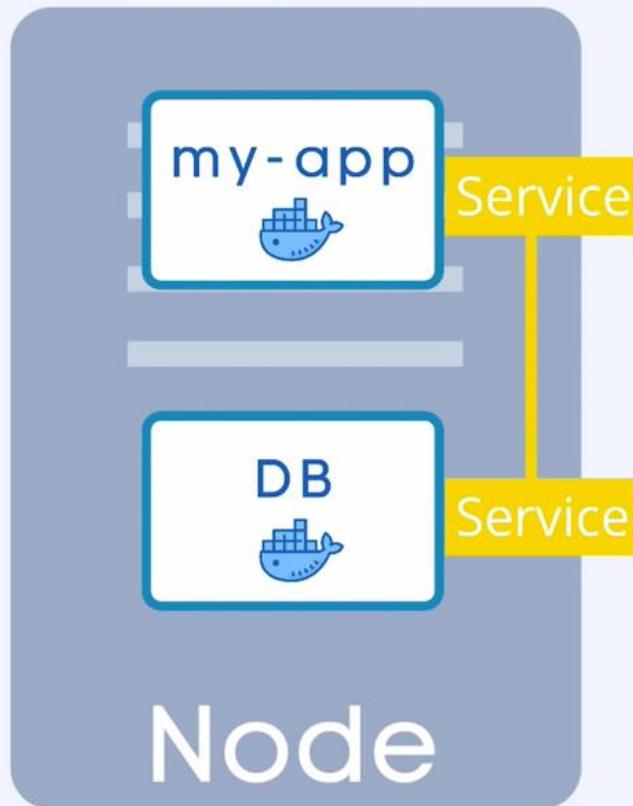


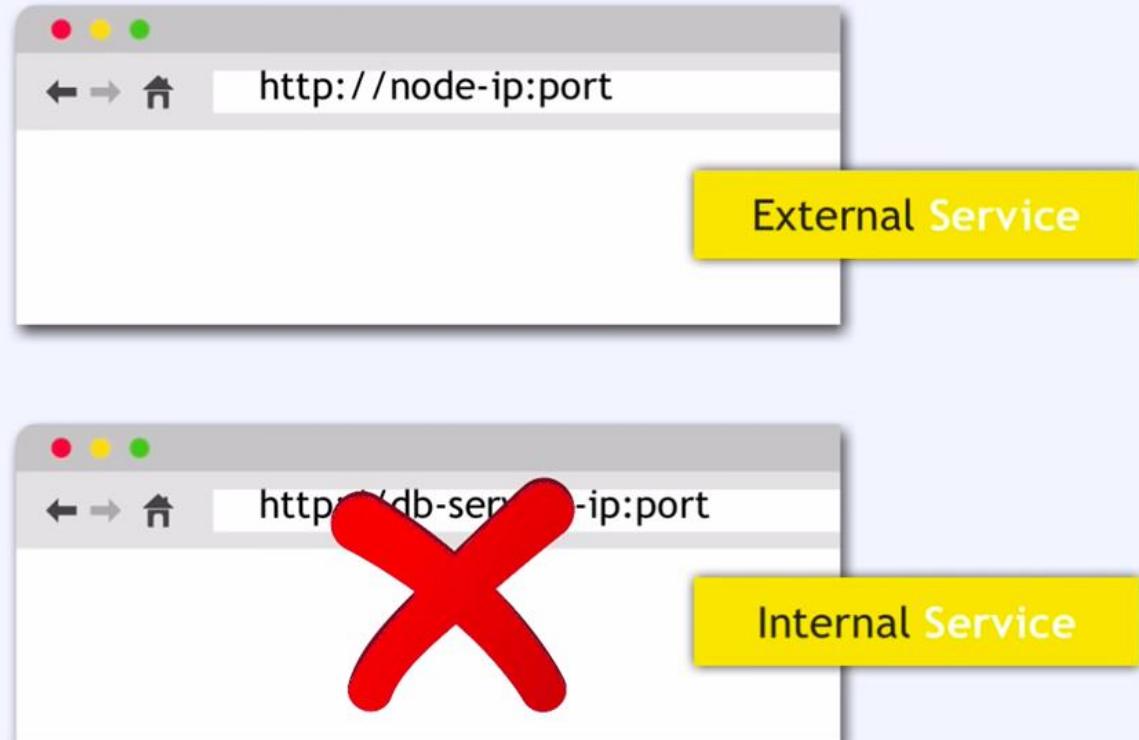
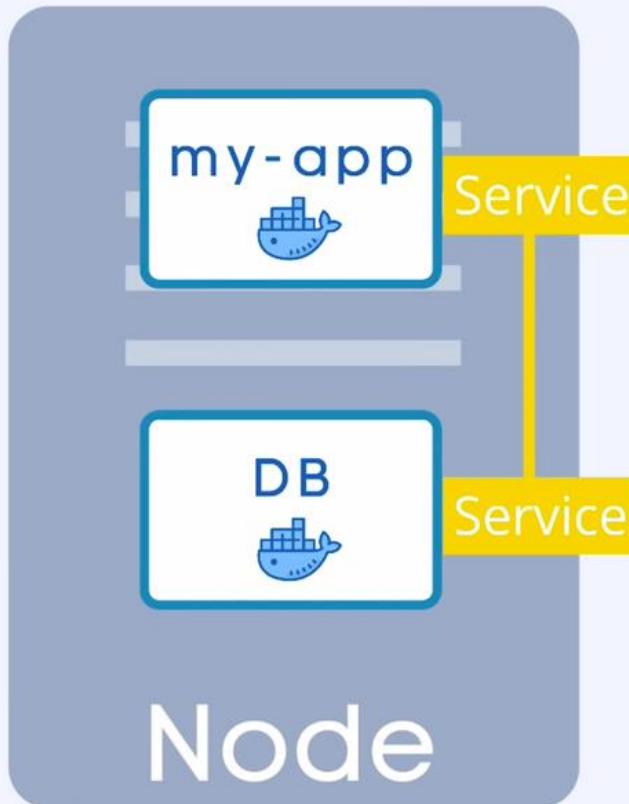


- ▶ Permanent IP address
- ▶ Lifecycle of Pod and Service  
not connected

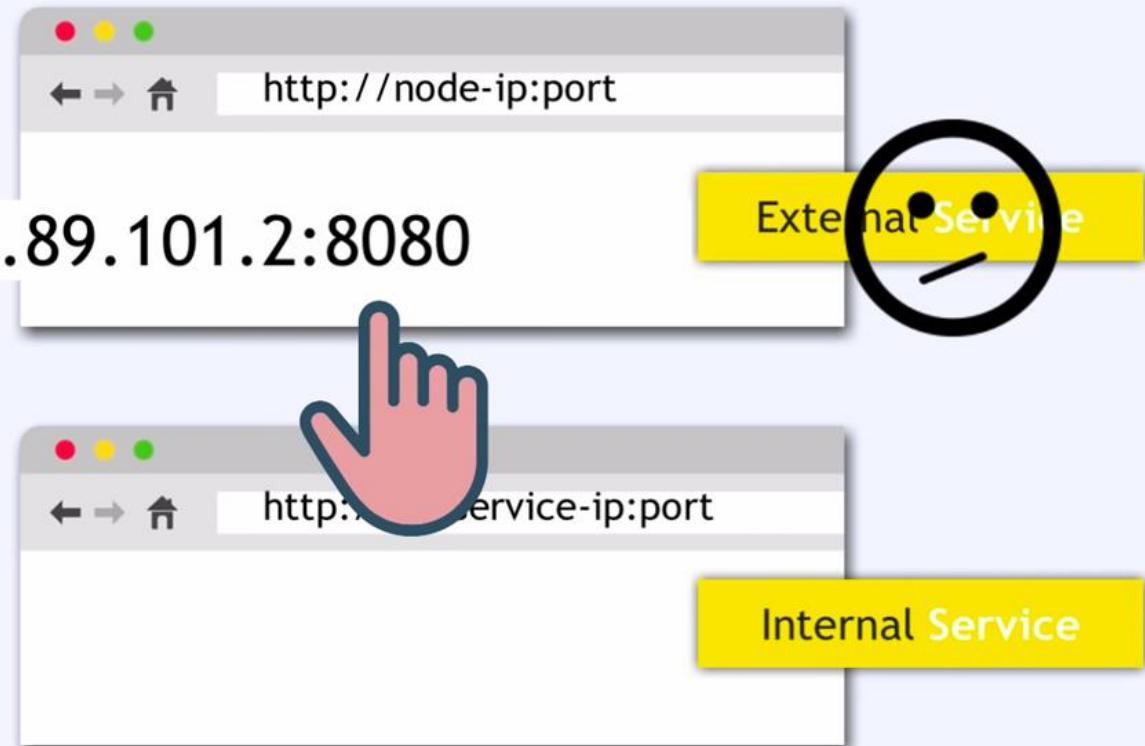
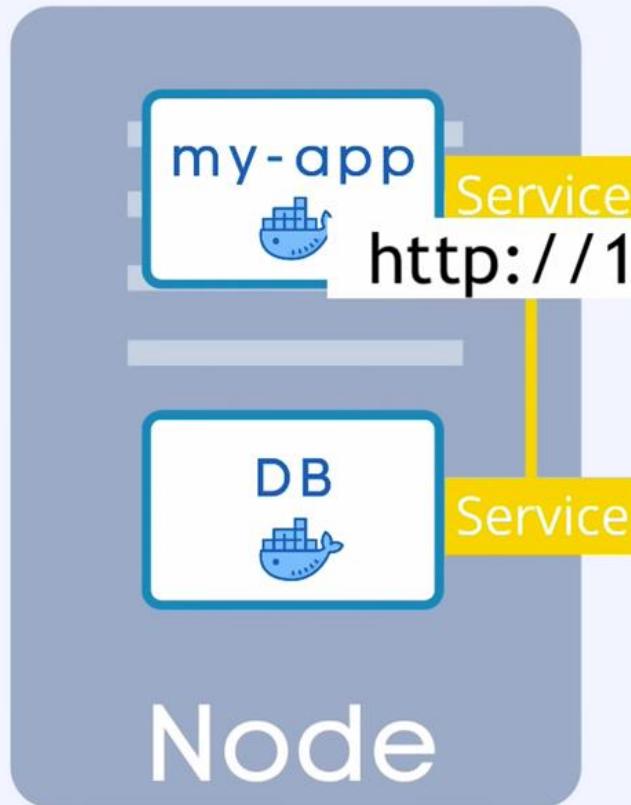


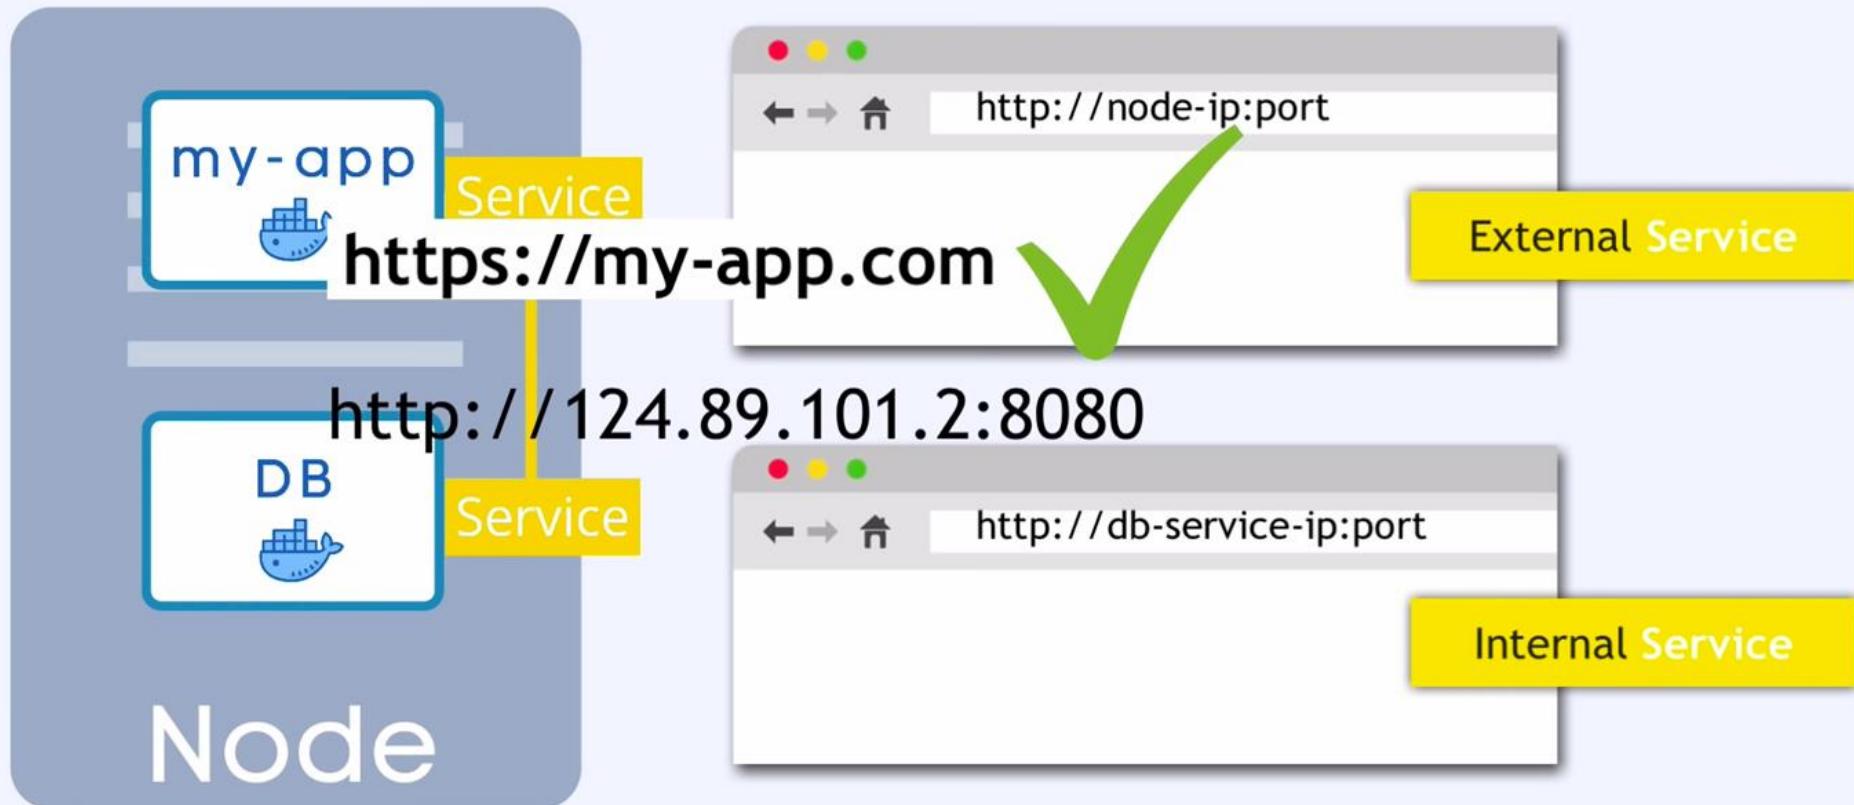


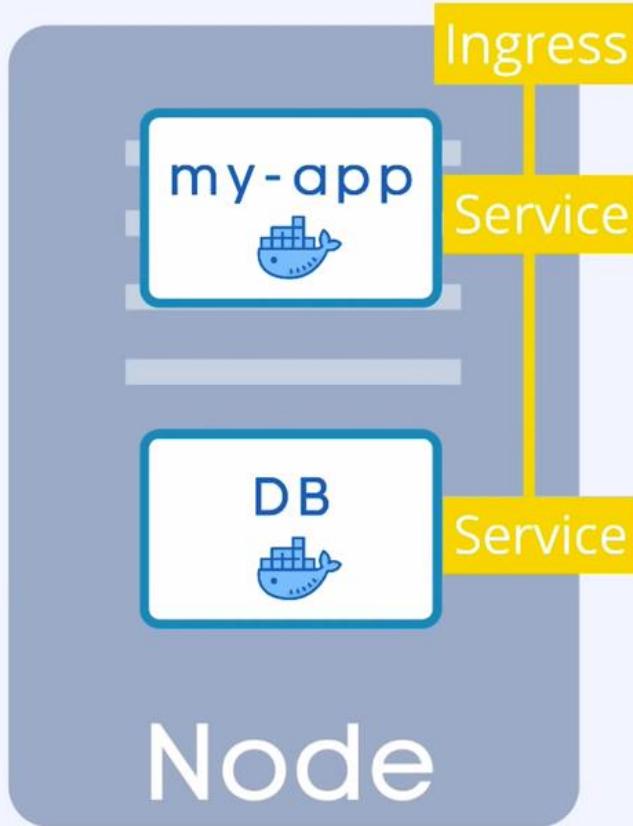




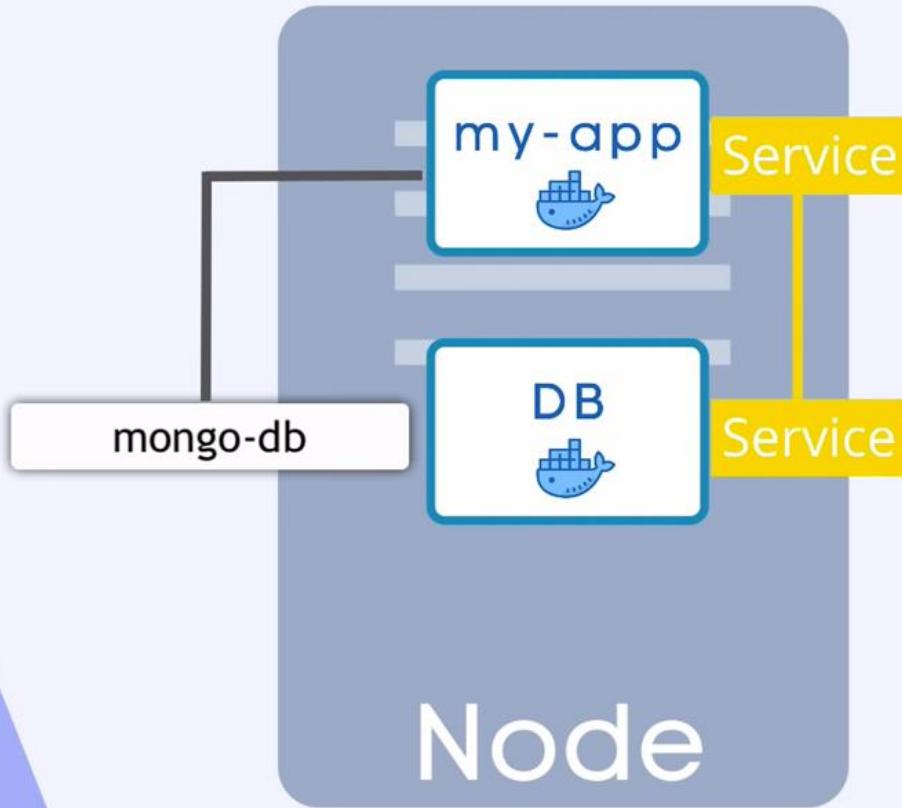
- ▶ You specify the **type of Service** on creation
- ▶ Internal Service is the **default type**





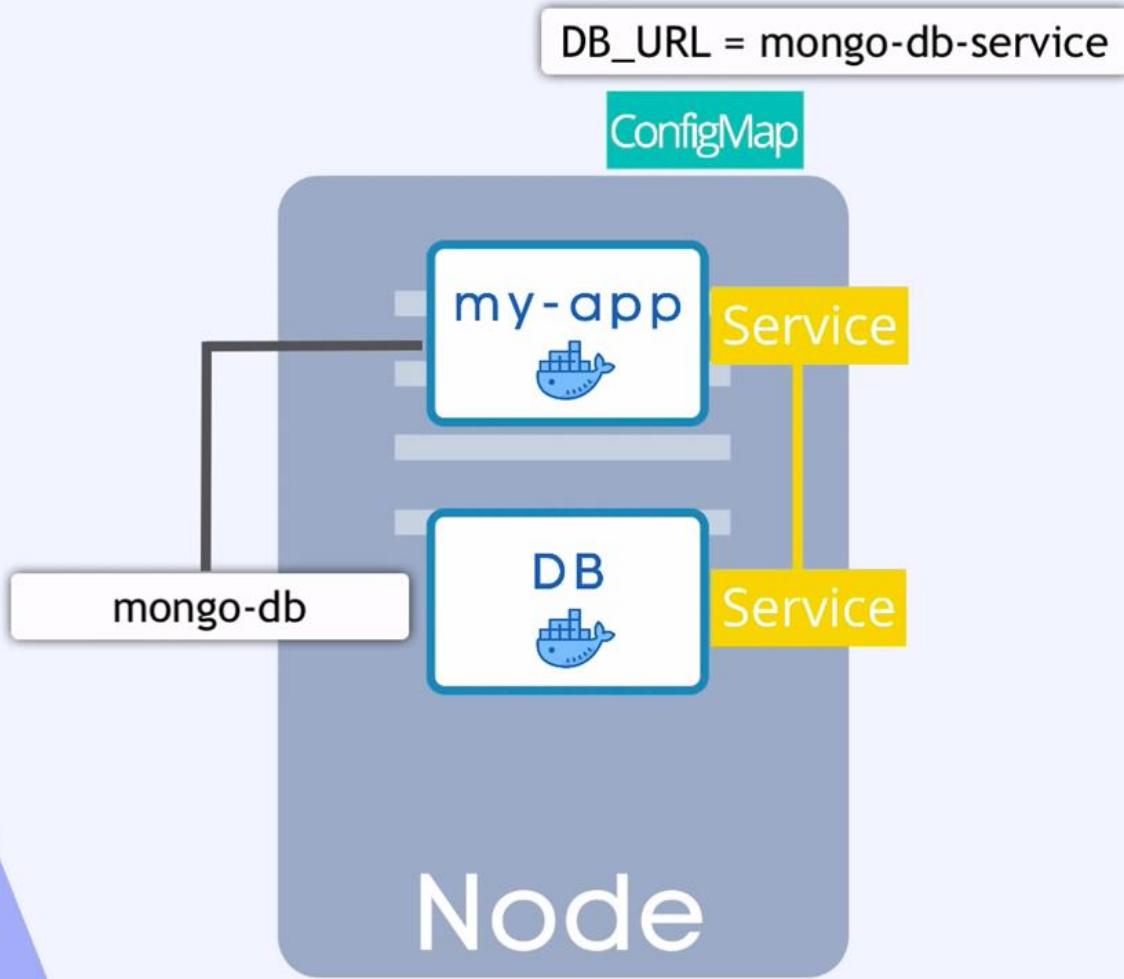


Database URL usually in the  
**built** application!

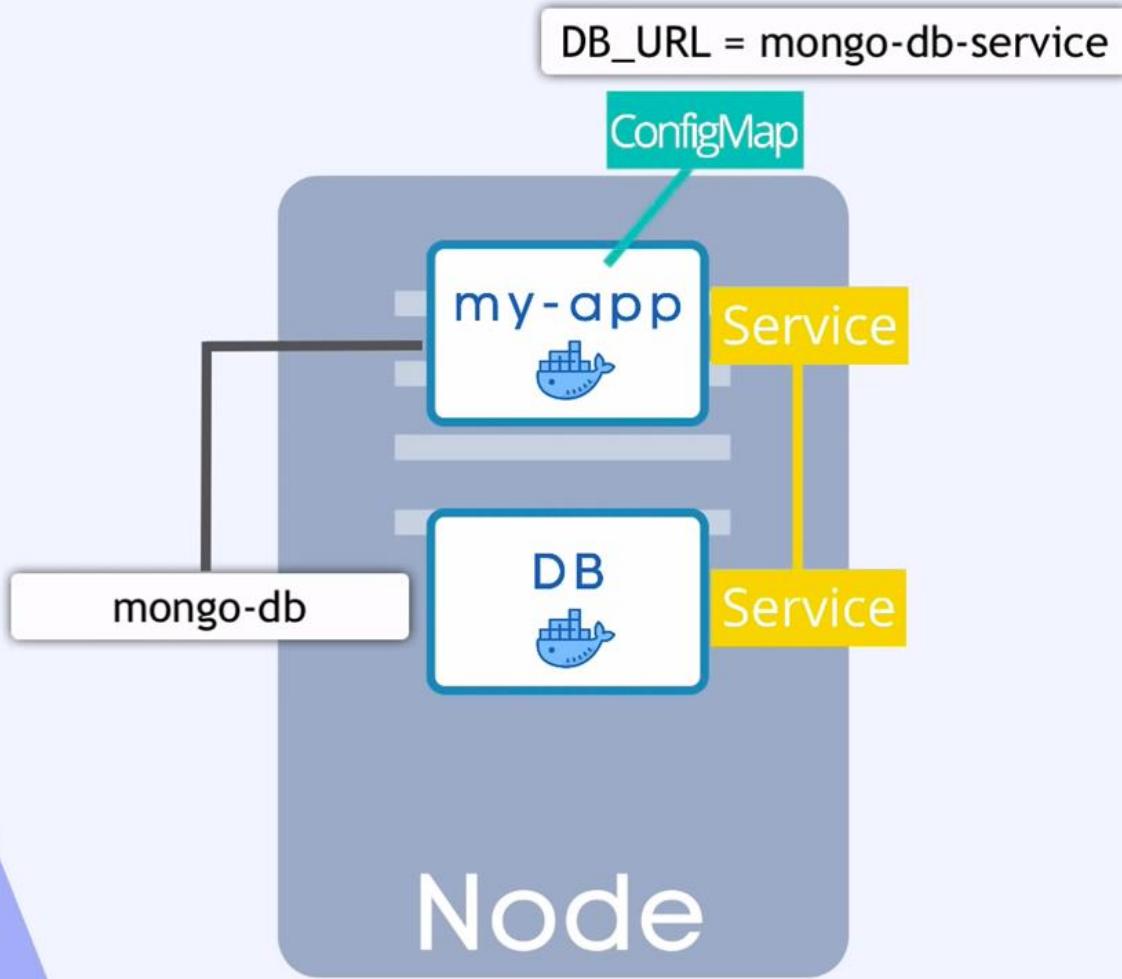


- Re-build
- Push it to repo
- Pull it in your Pod

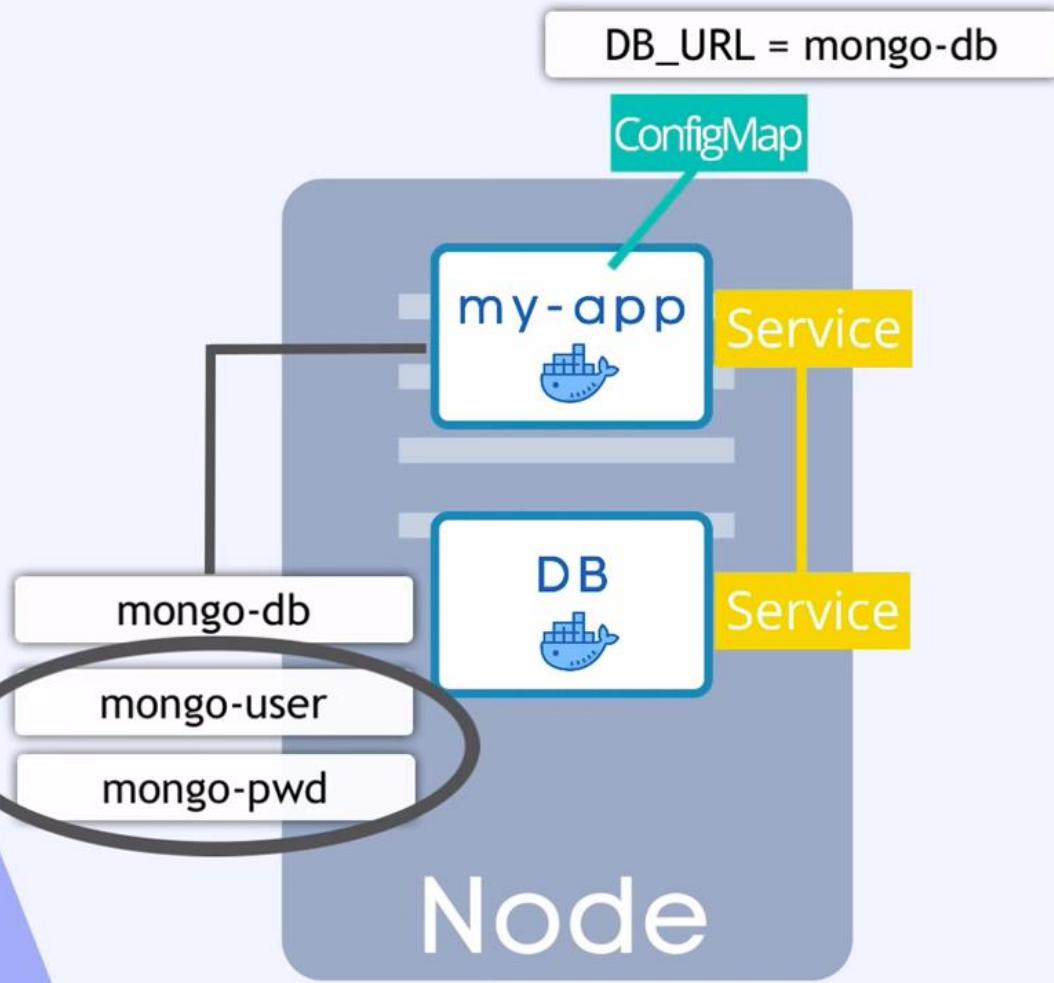


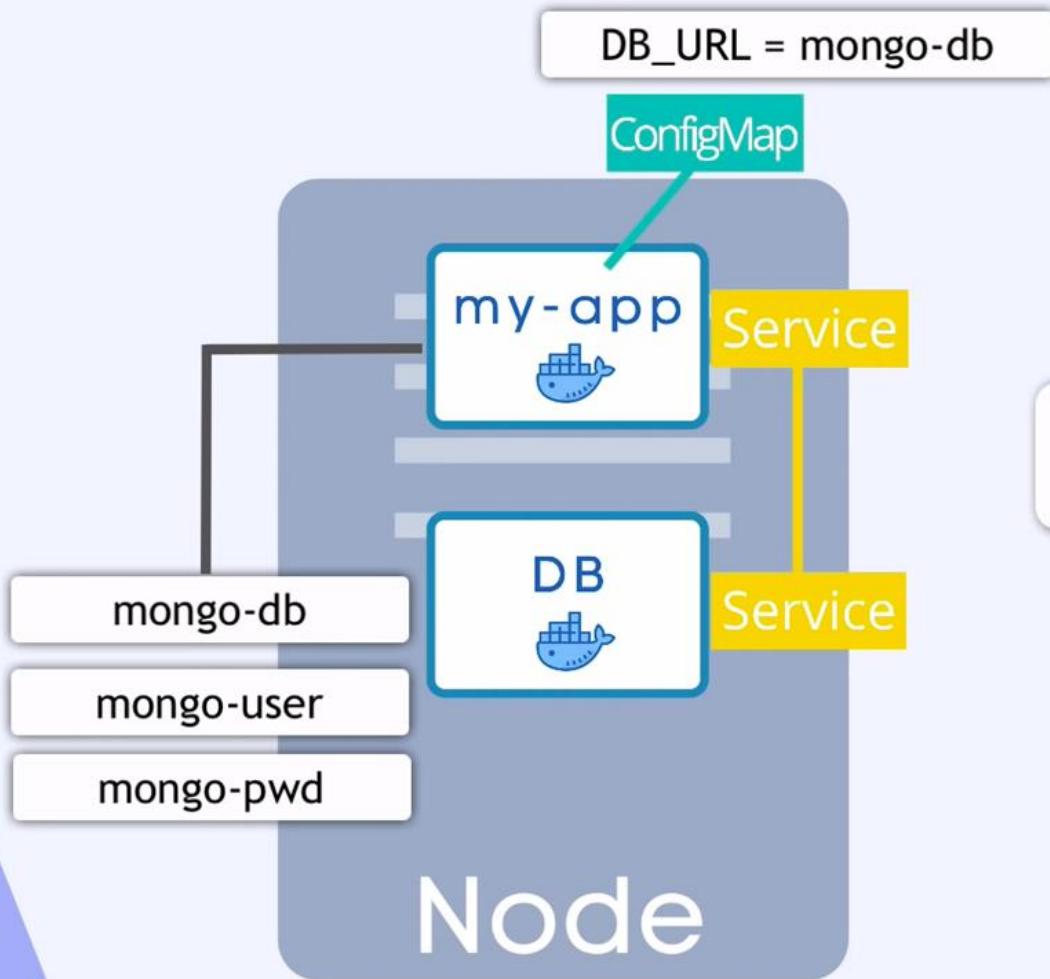


- ▶ **External Configuration** of your application

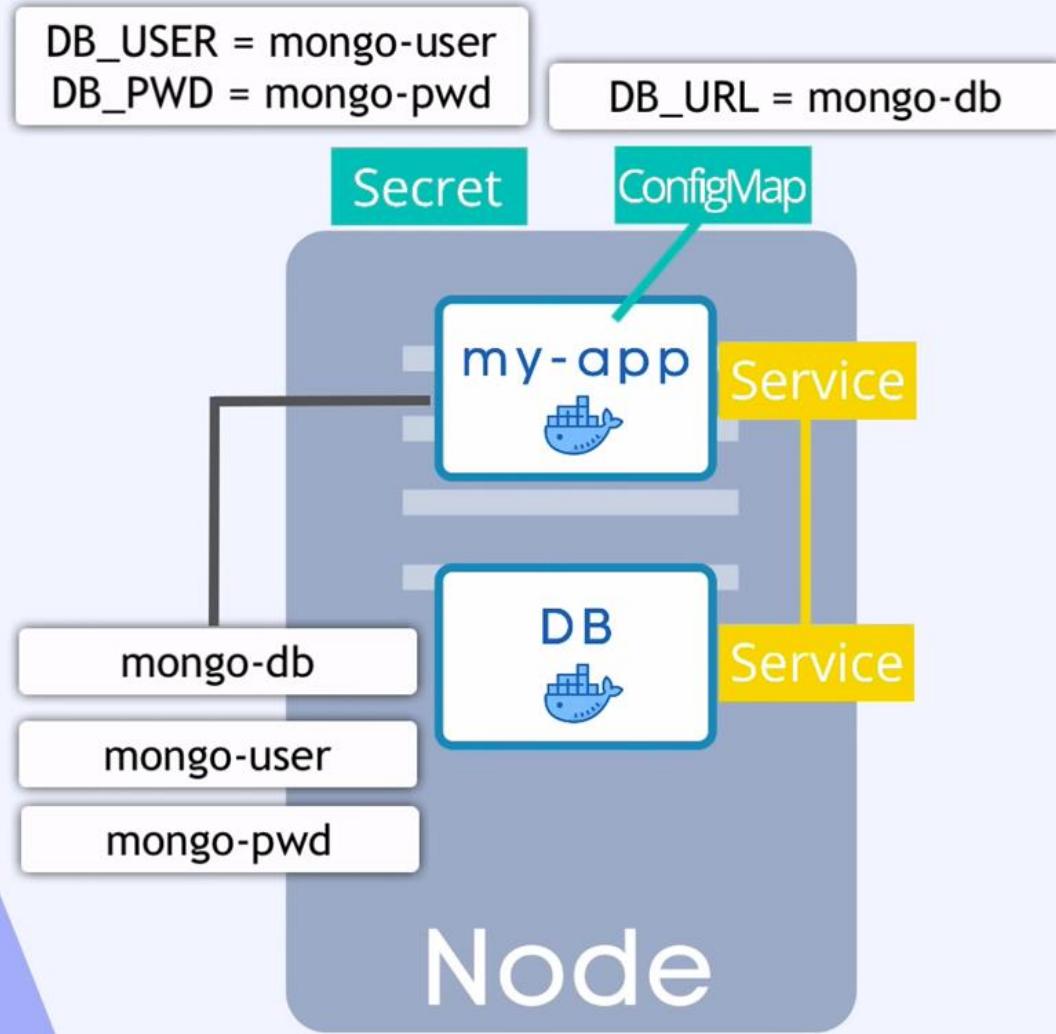


- ▶ **External Configuration** of your application





! ConfigMap is for non-confidential data only!



► Used to store **secret** data

DB\_USER = mongo-user  
DB\_PWD = mongo-pwd

DB\_URL = mongo-db

Secret

mongo-db

mongo-user

mongo-pwd

! The built-in security mechanism is  
not enabled by default!

#### Caution:

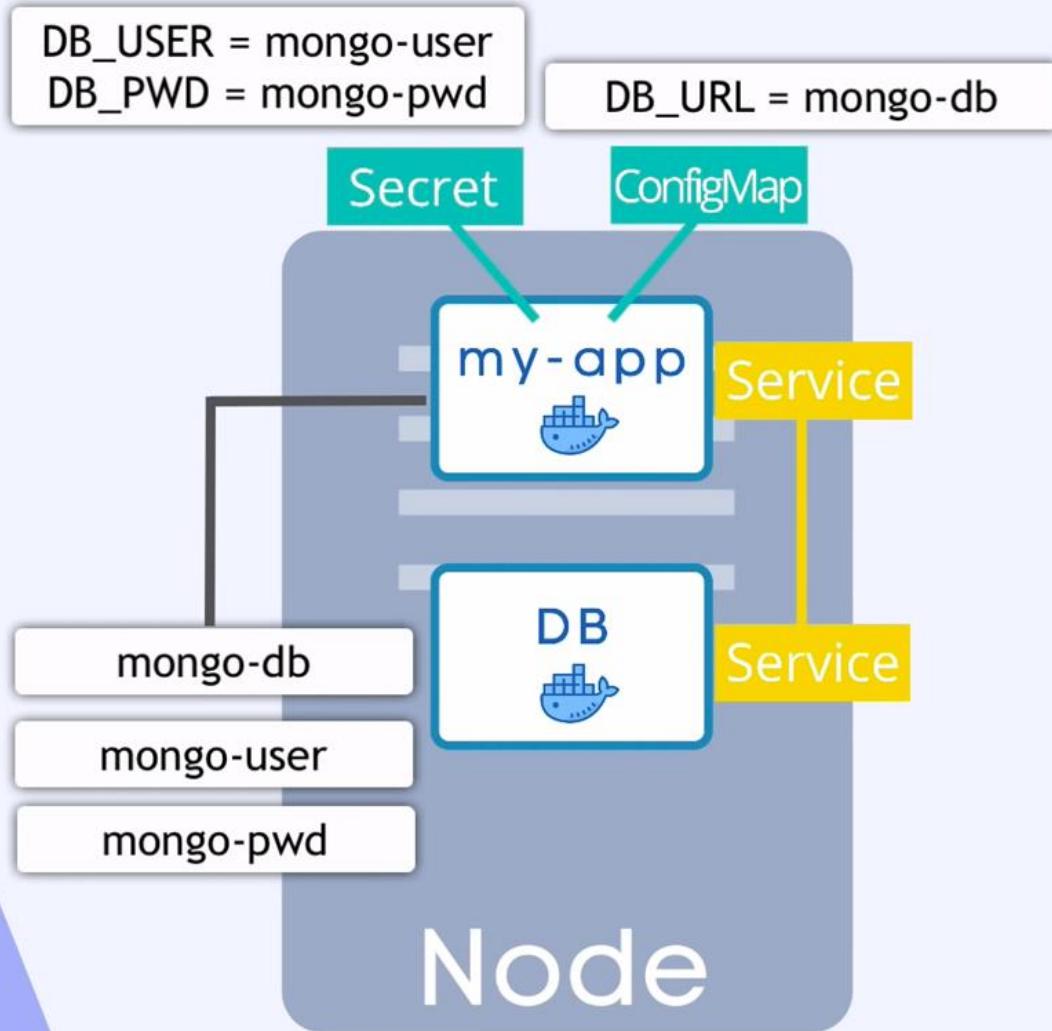
Kubernetes Secrets are, by default, stored unencrypted in the API server's underlying data store (etcd). Anyone with API access can retrieve or modify a Secret, and so can anyone with access to etcd. Additionally, anyone who is authorized to create a Pod in a namespace can use that access to read any Secret in that namespace; this includes indirect access such as the ability to create a Deployment.

In order to safely use Secrets, take at least the following steps:

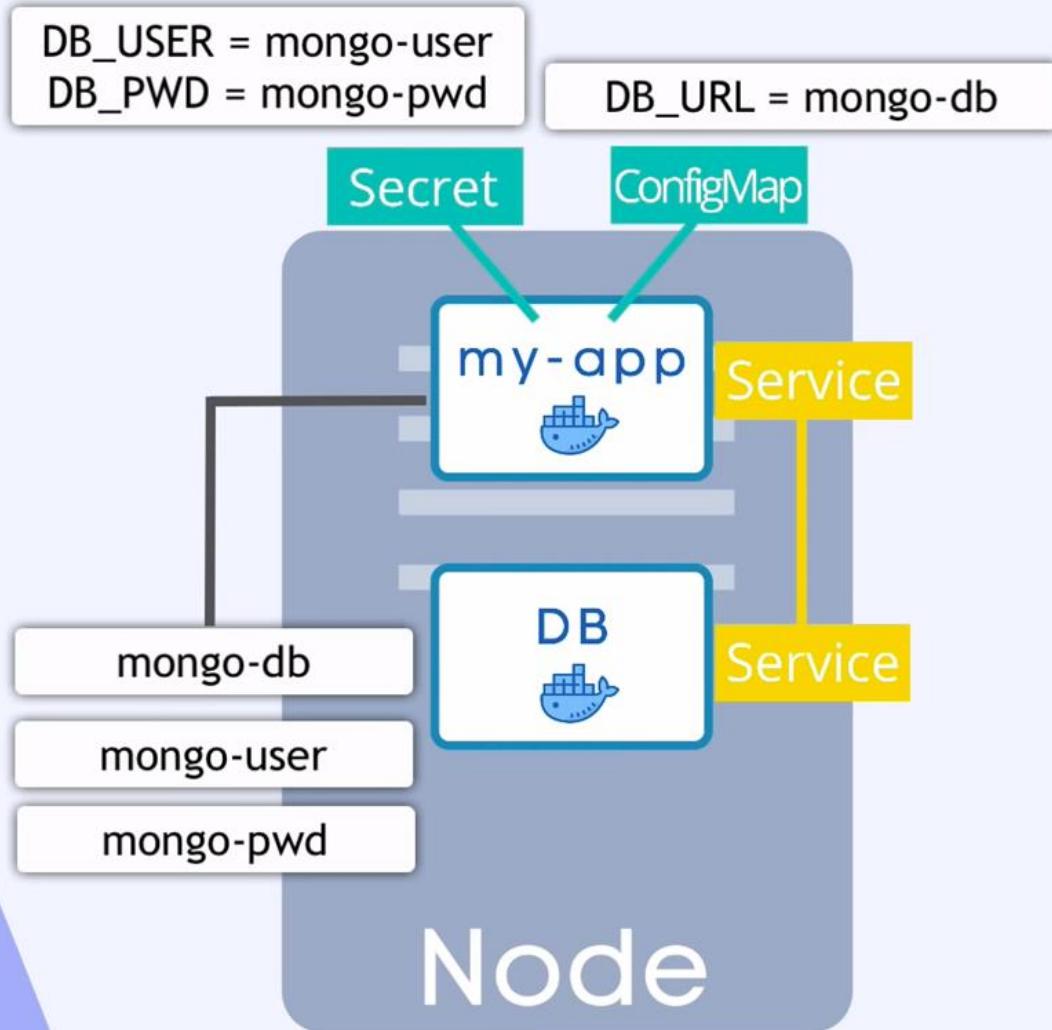
1. [Enable Encryption at Rest](#) for Secrets.
2. Enable or configure [RBAC rules](#) that restrict reading data in Secrets (including via indirect means).
3. Where appropriate, also use mechanisms such as RBAC to limit which principals are allowed to create new Secrets or replace existing ones.



Check out the official docs  
for more info on this:

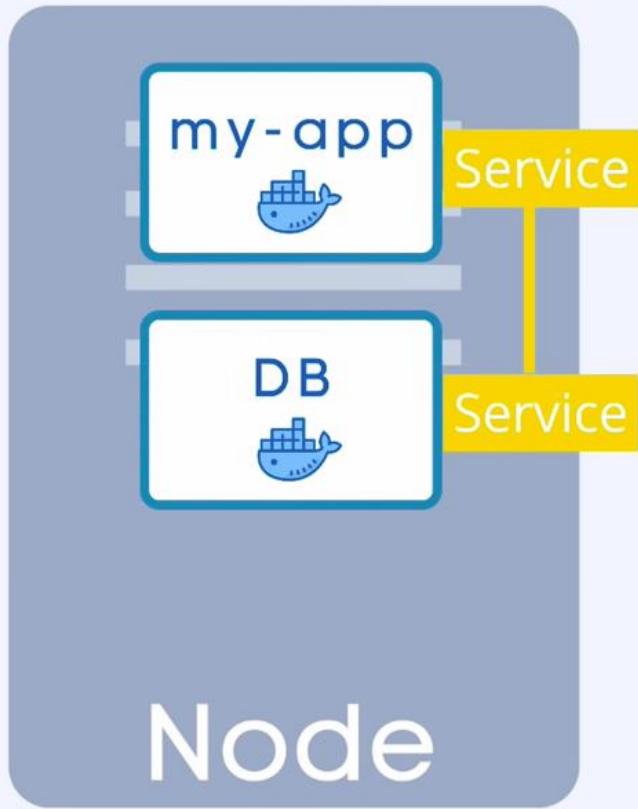


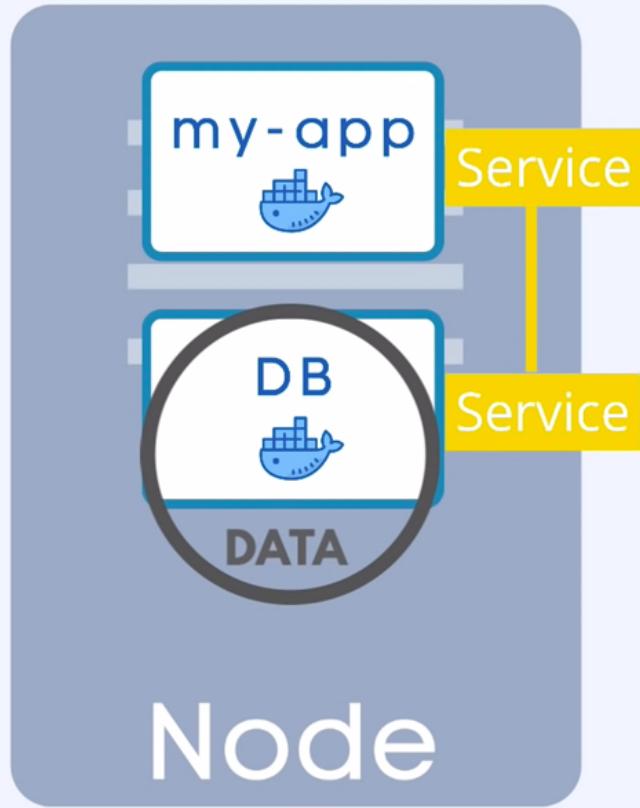
- ▶ Used to store **secret** data
- ▶ Reference Secret in Deployment/Pod

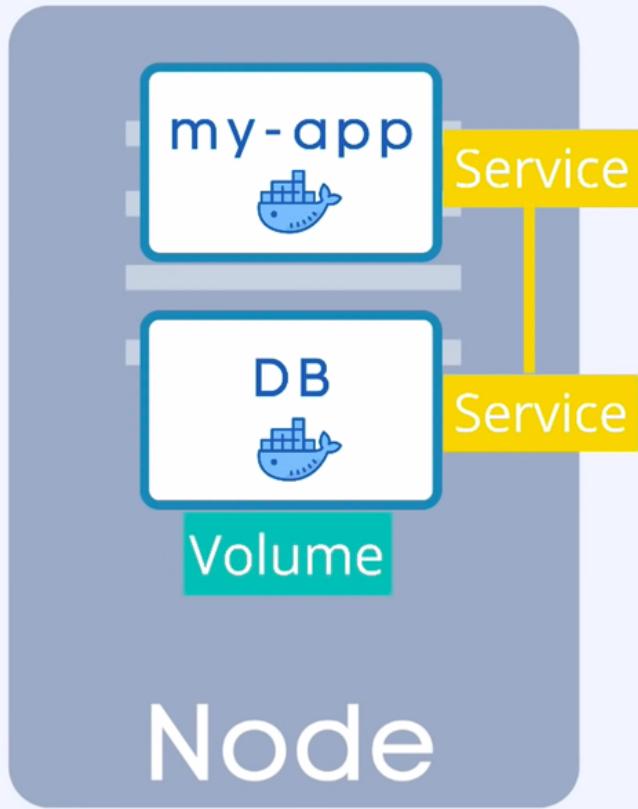


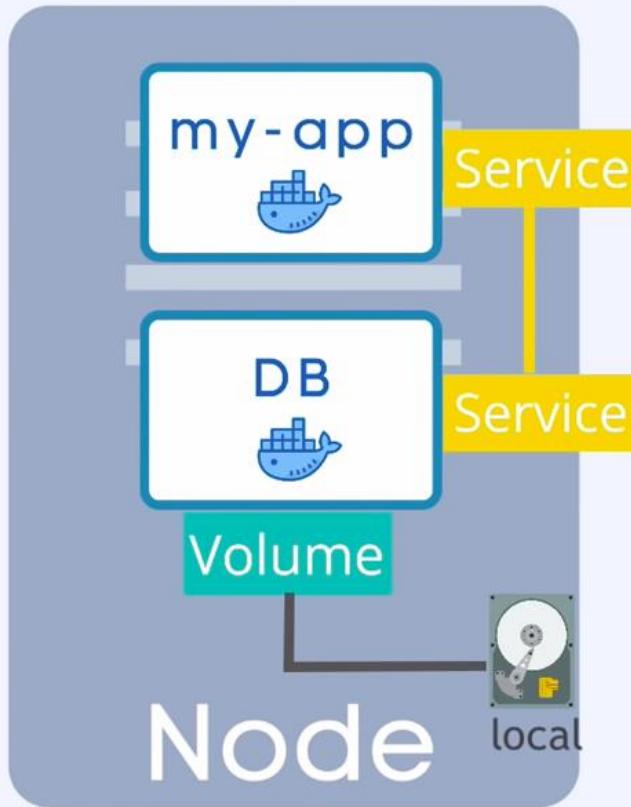
- ▶ Used to store **secret** data
- ▶ Reference Secret in Deployment/Pod

Use it as environment variables or as a properties file

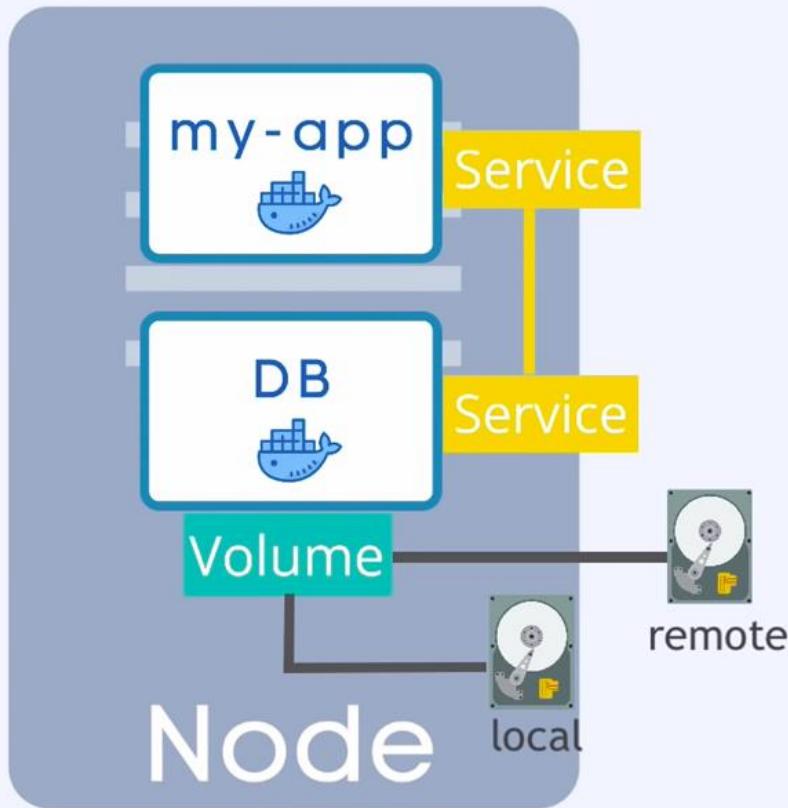




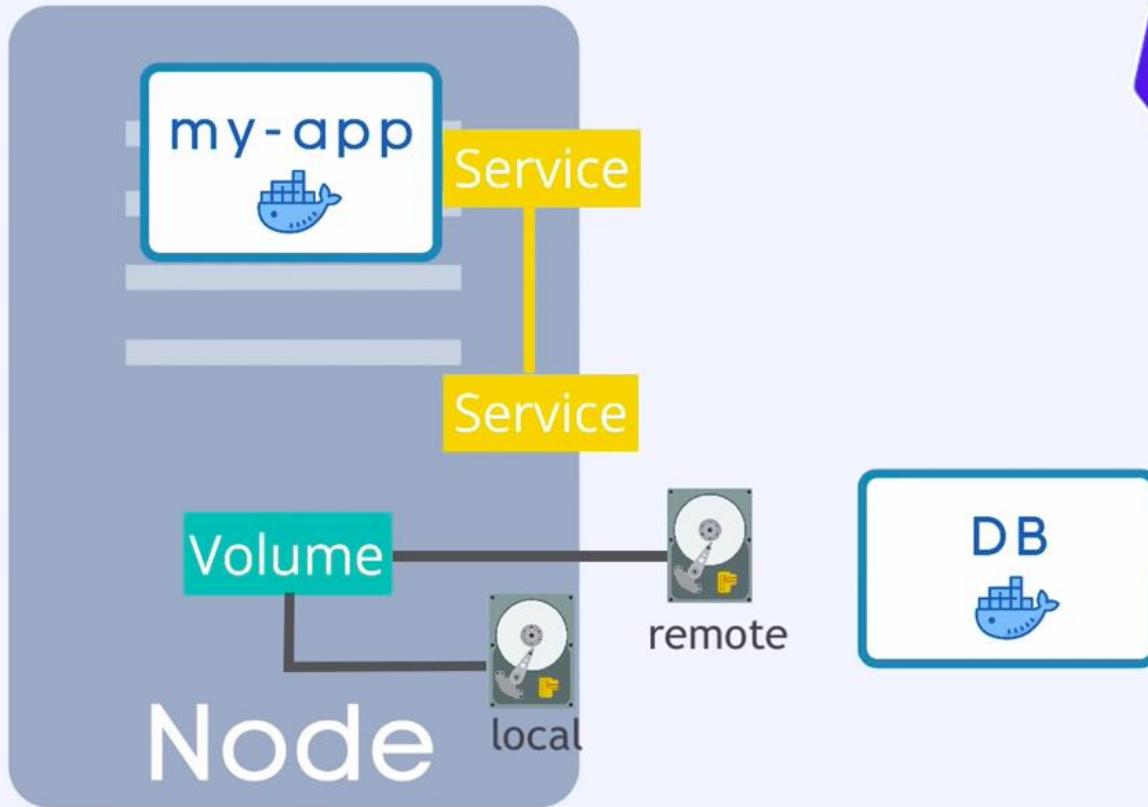


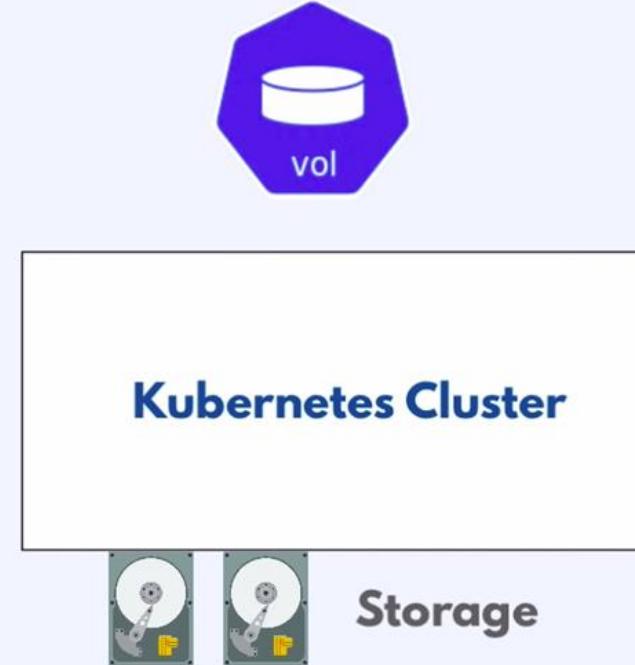
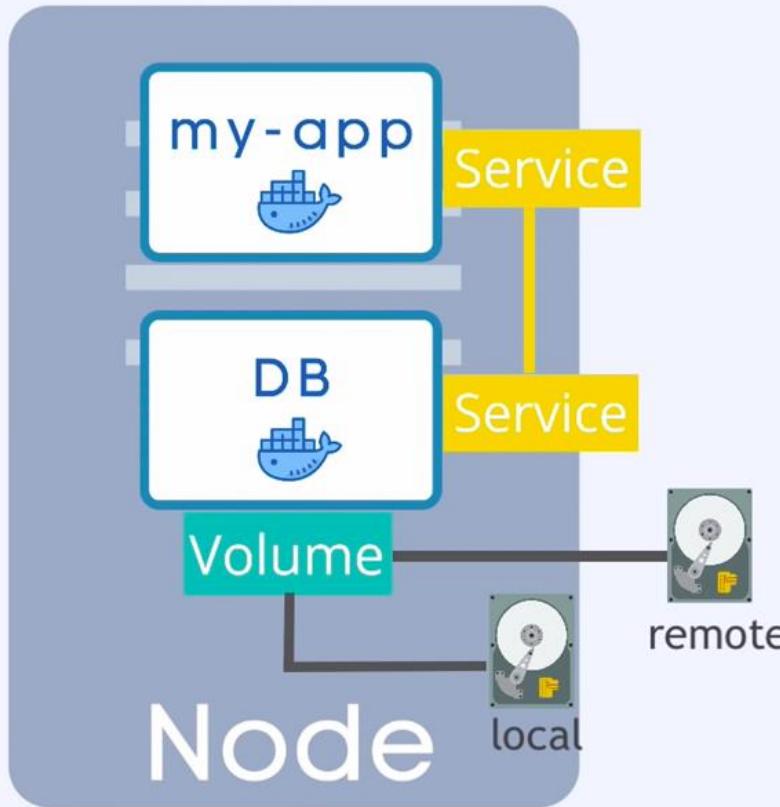


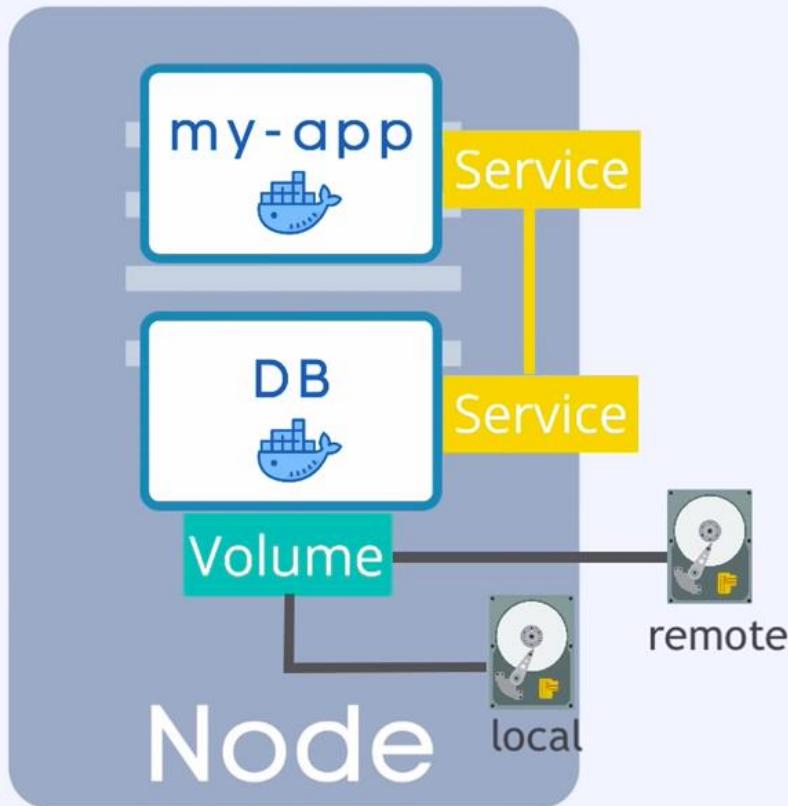
► Storage on **local** machine



- ▶ Storage on **local** machine
- ▶ Or **remote**, outside of the K8s cluster



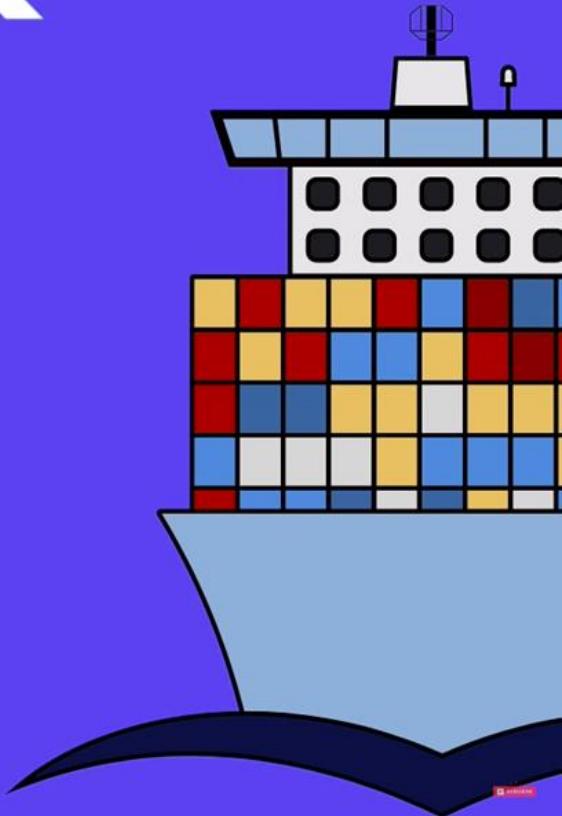


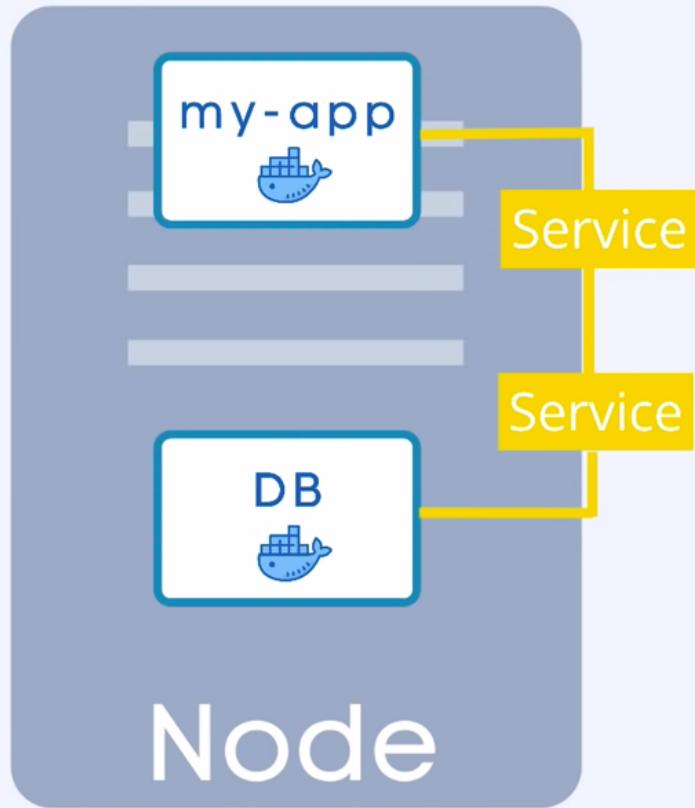


! Kubernetes doesn't manage data persistence!



# Deployment & StatefulSet

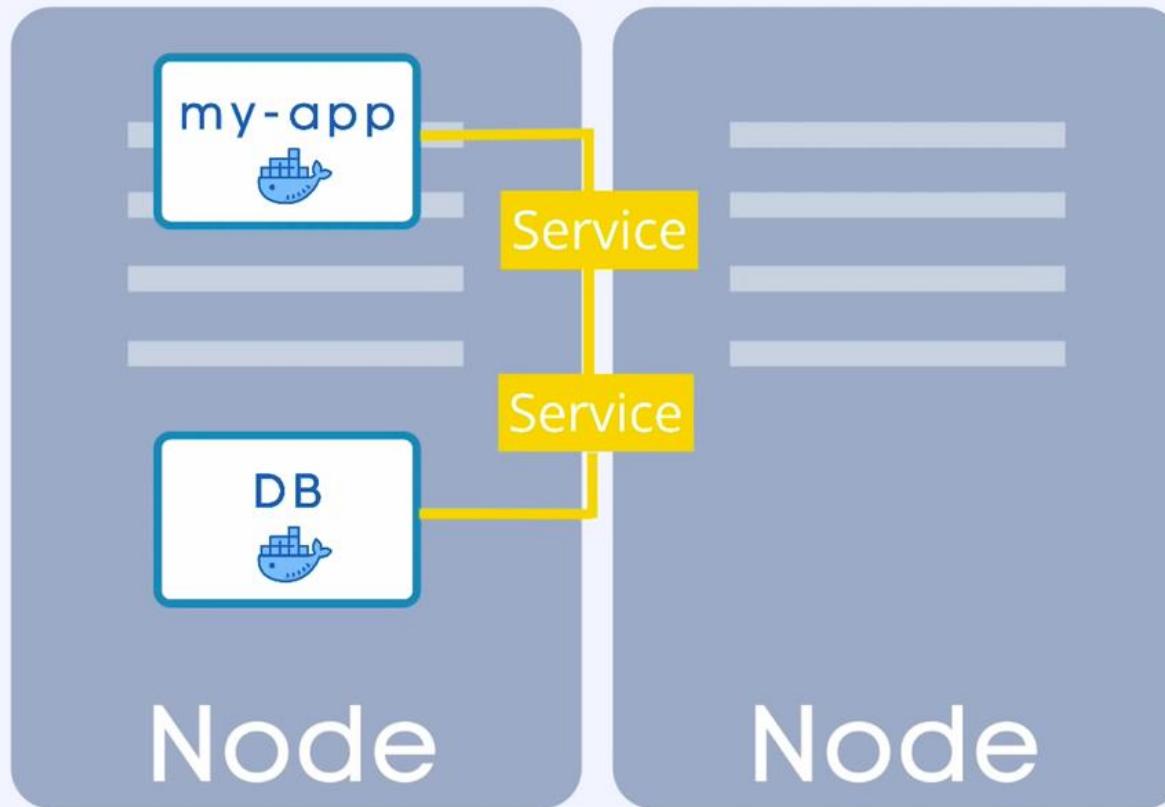


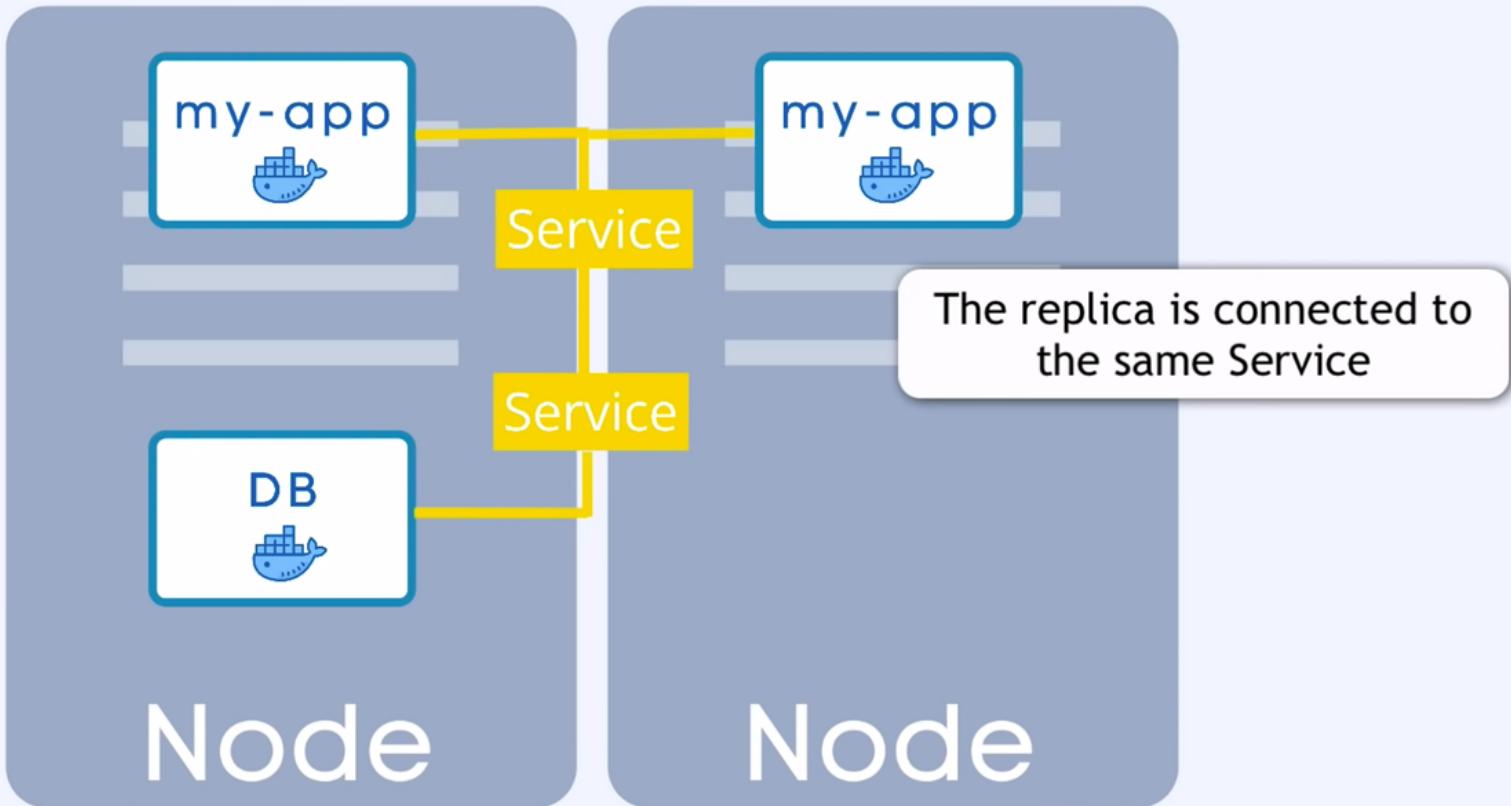


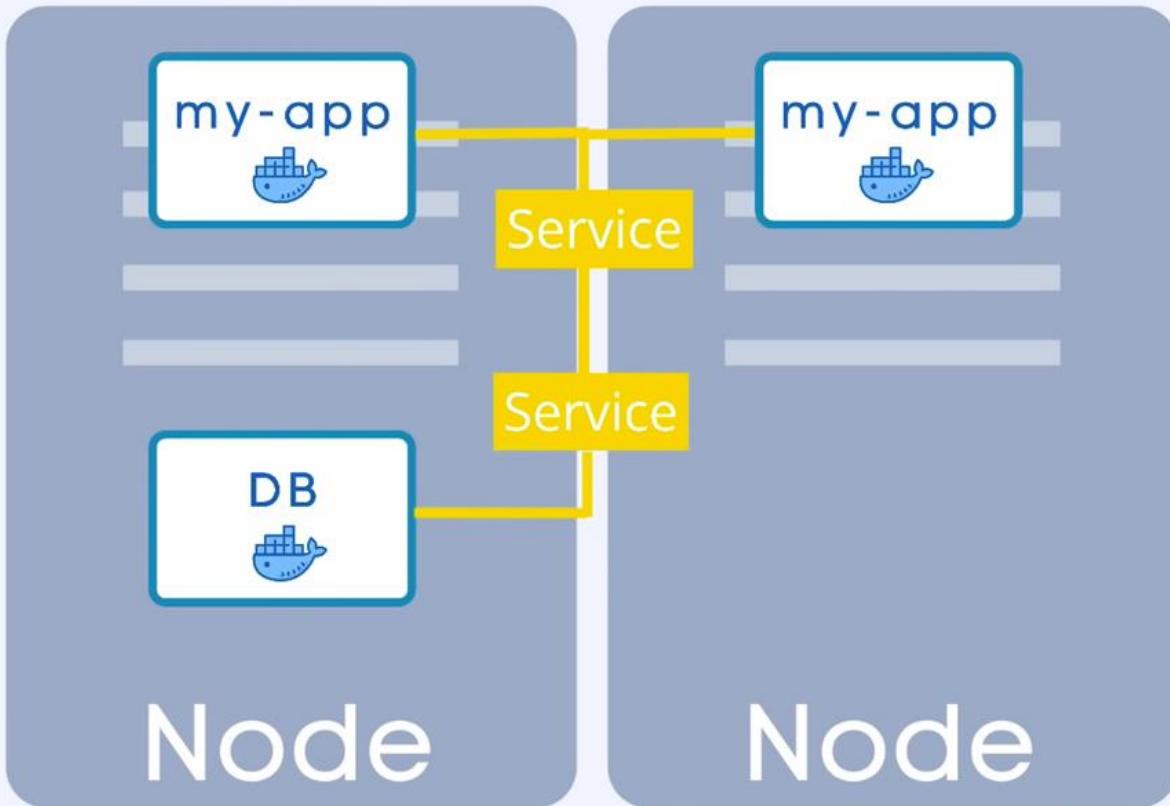
**Distributed Systems**

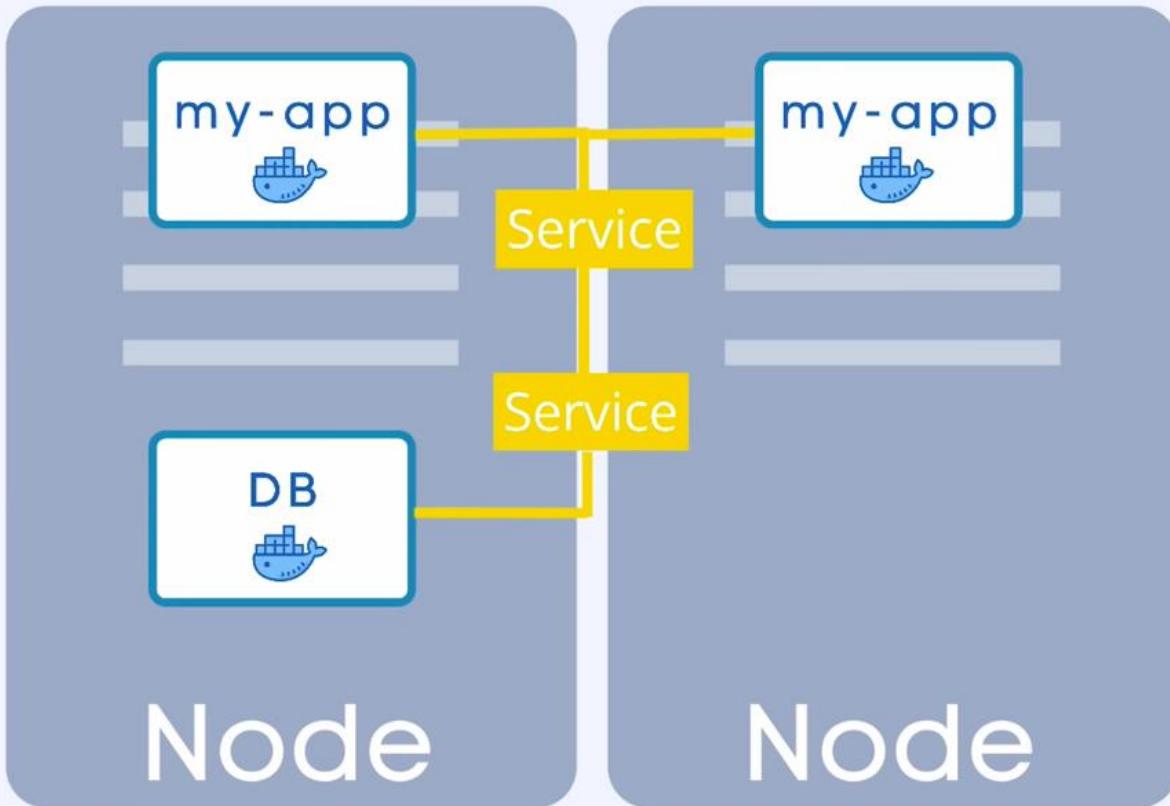


Replicate everything

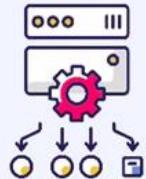


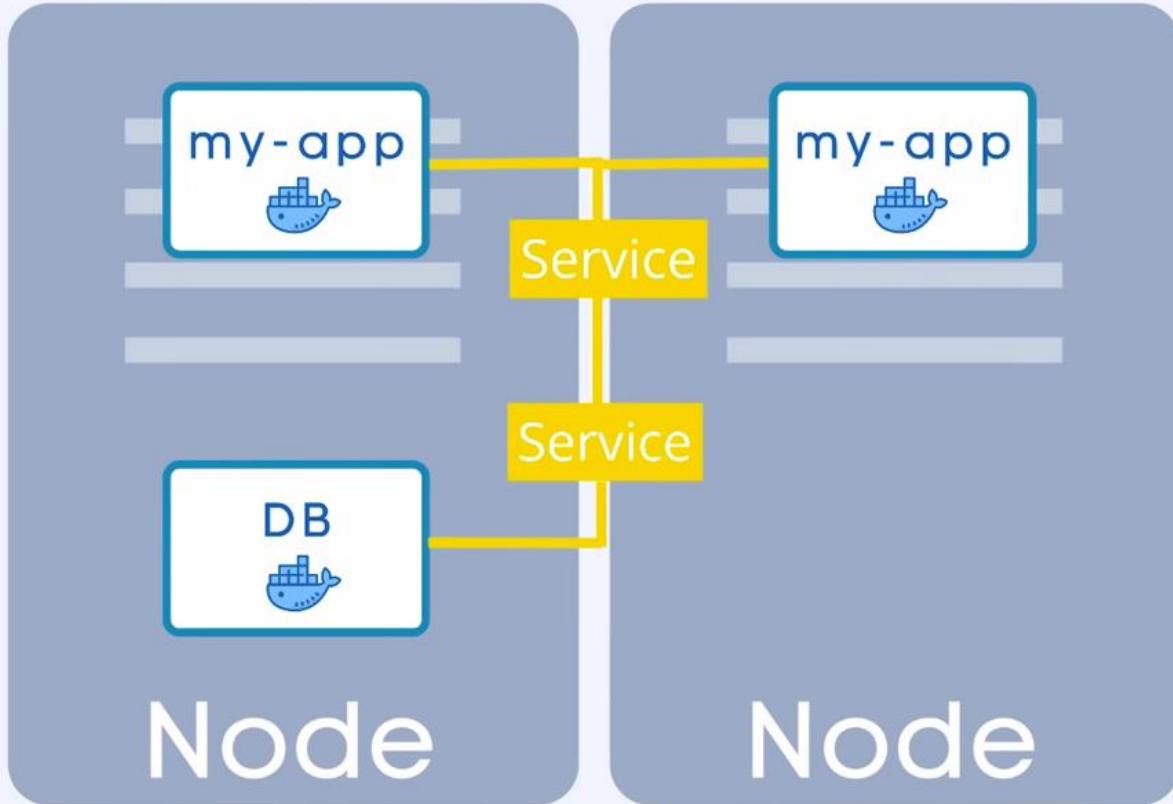




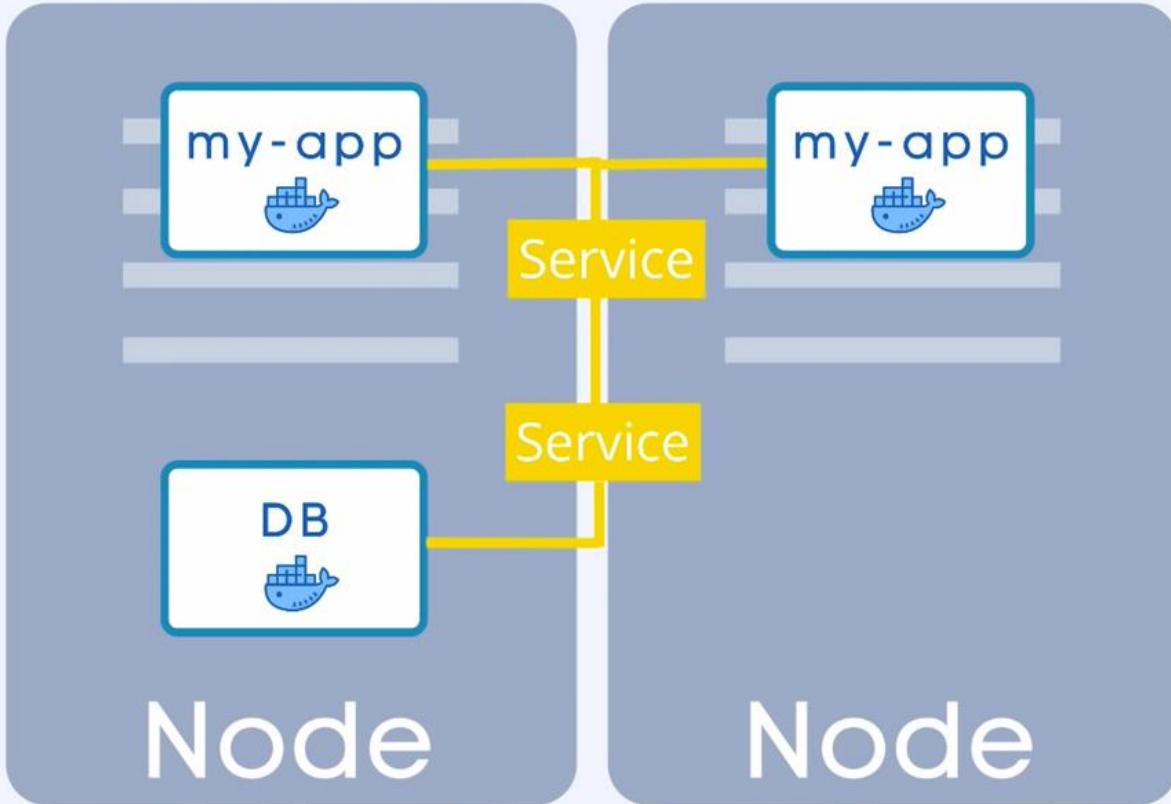


- ▶ permanent IP
- ▶ load balancer



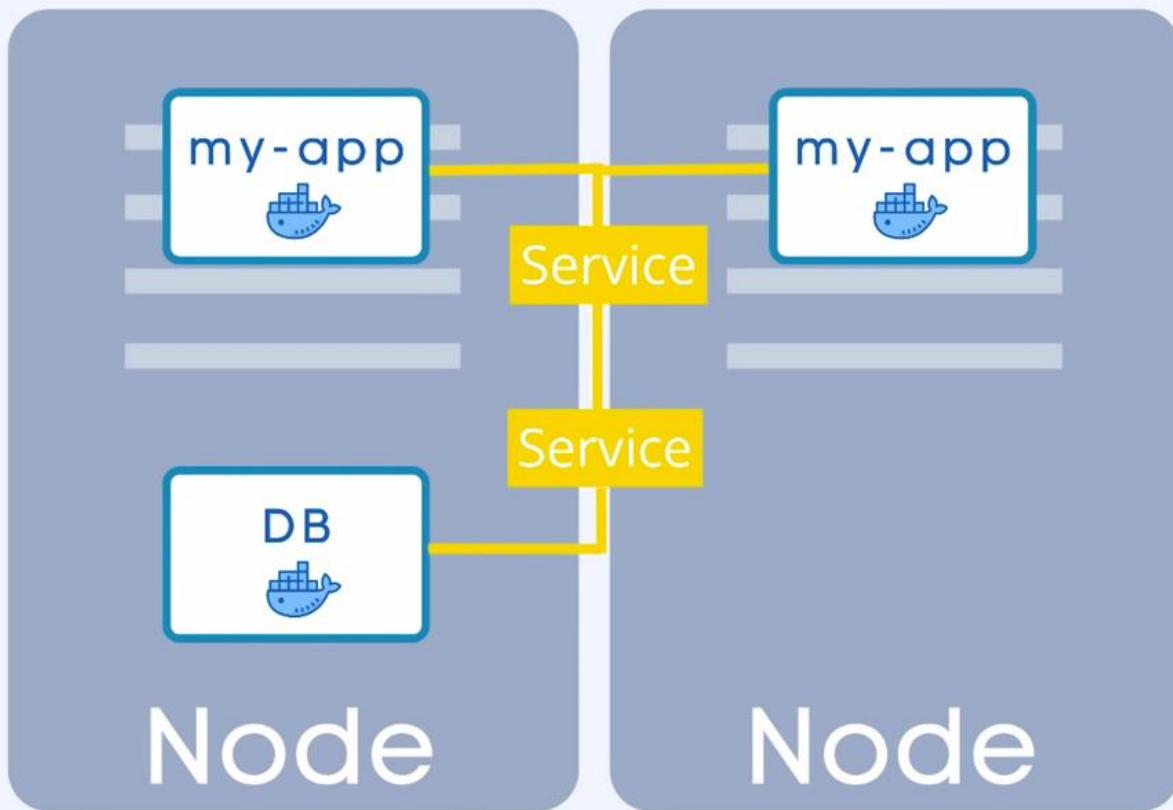


Define blueprint for Pods

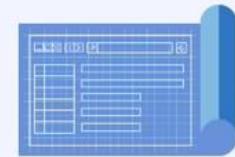


Define **blueprint** for Pods

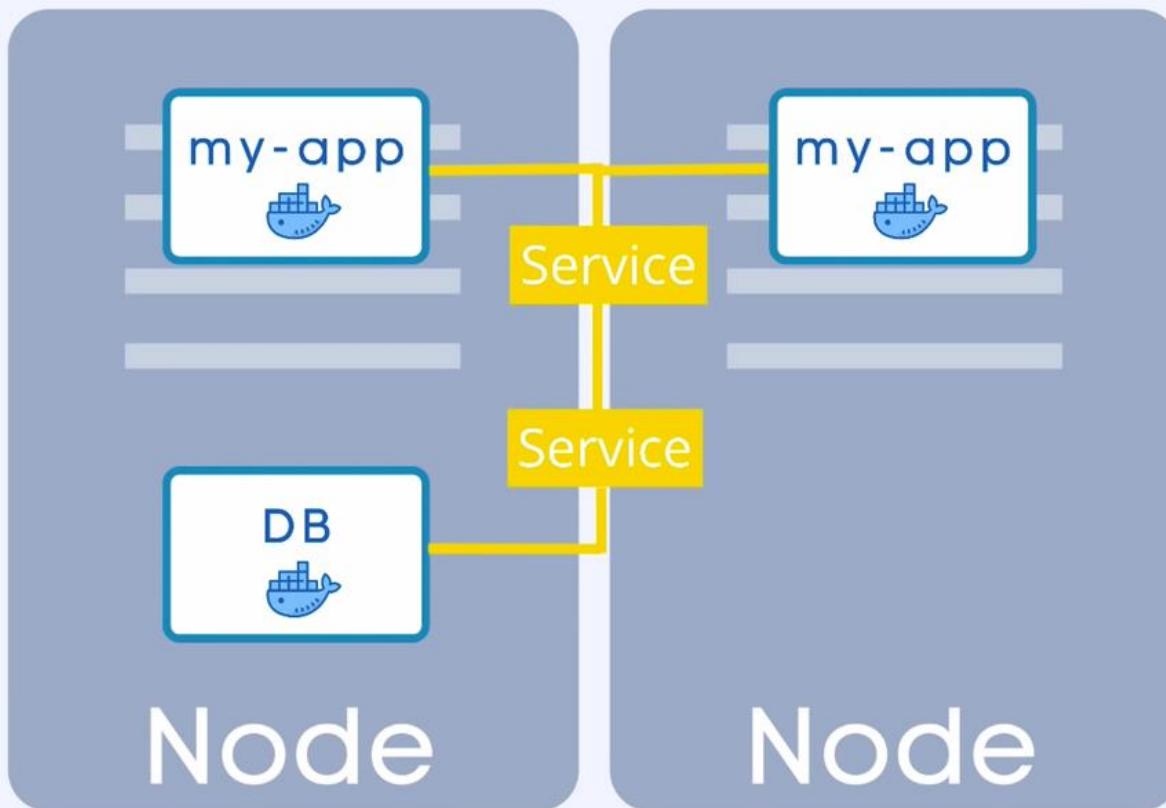
- ▶ Specify how many replicas you want to have



## DEPLOYMENT



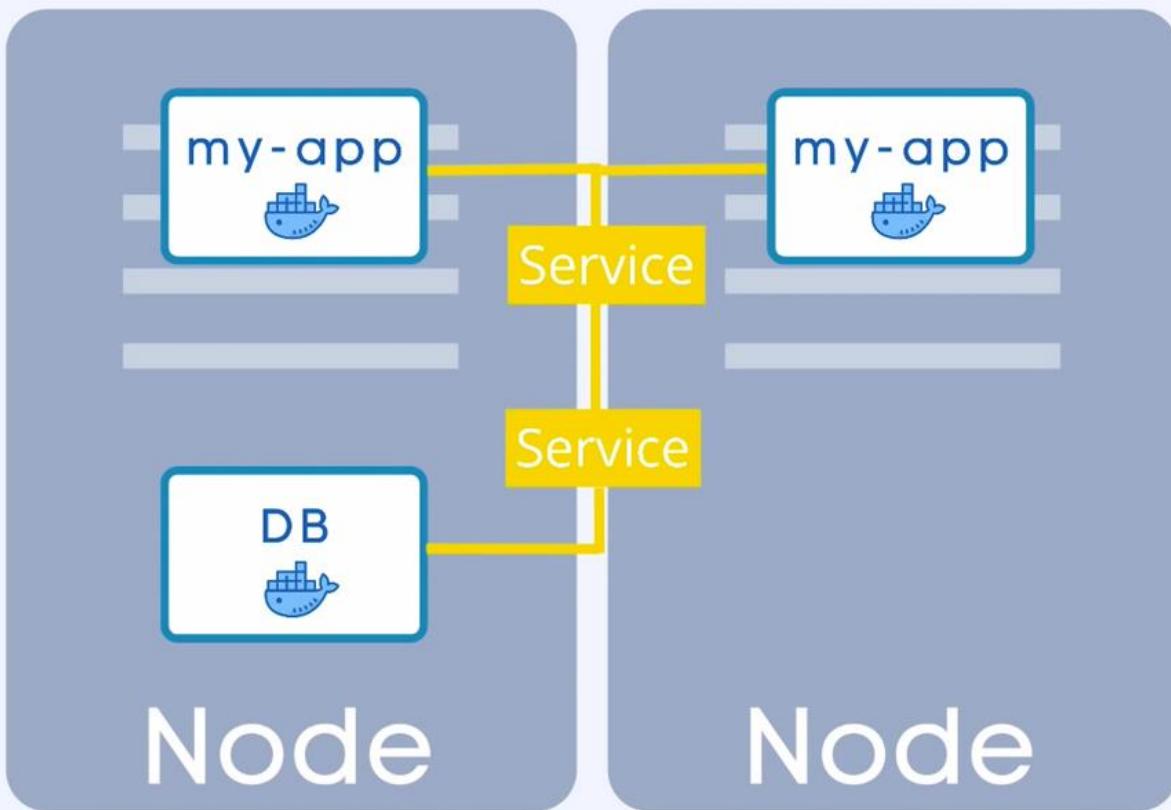
- ▶ Blueprint for "my-app" Pods



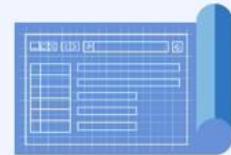
## DEPLOYMENT



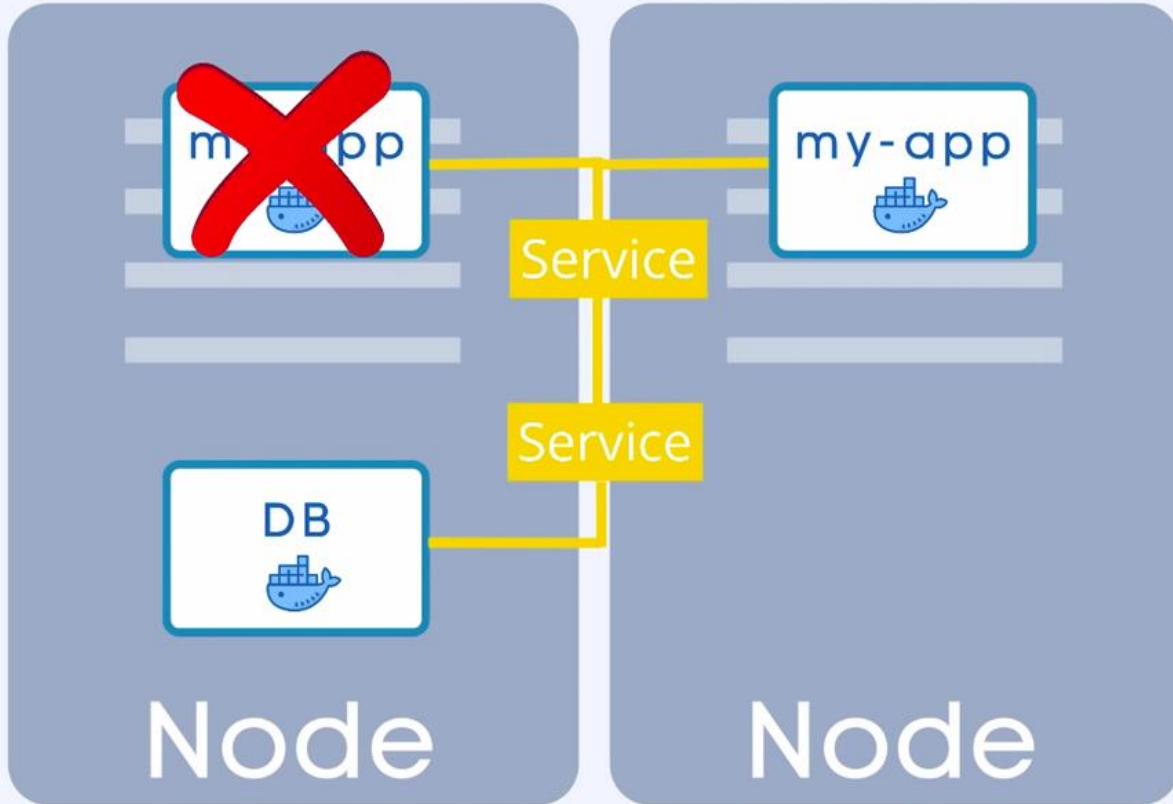
- ▶ Blueprint for "my-app" Pods
- ▶ You create Deployments

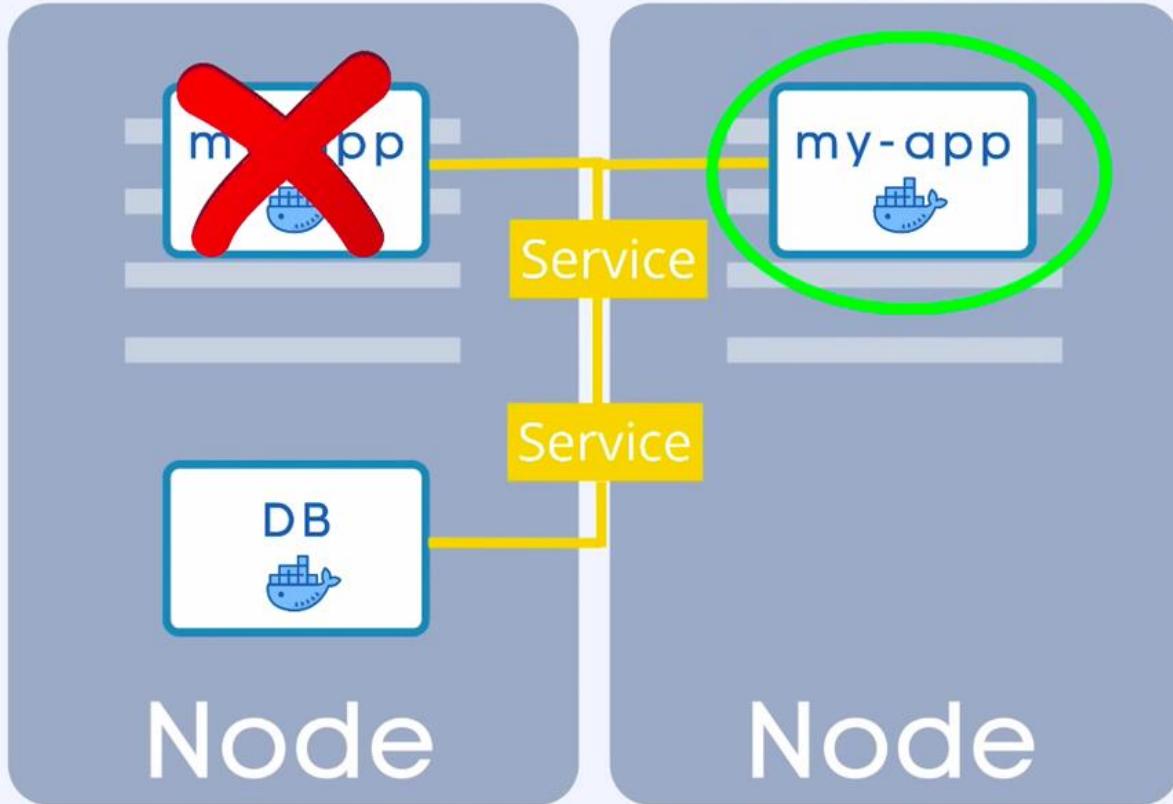


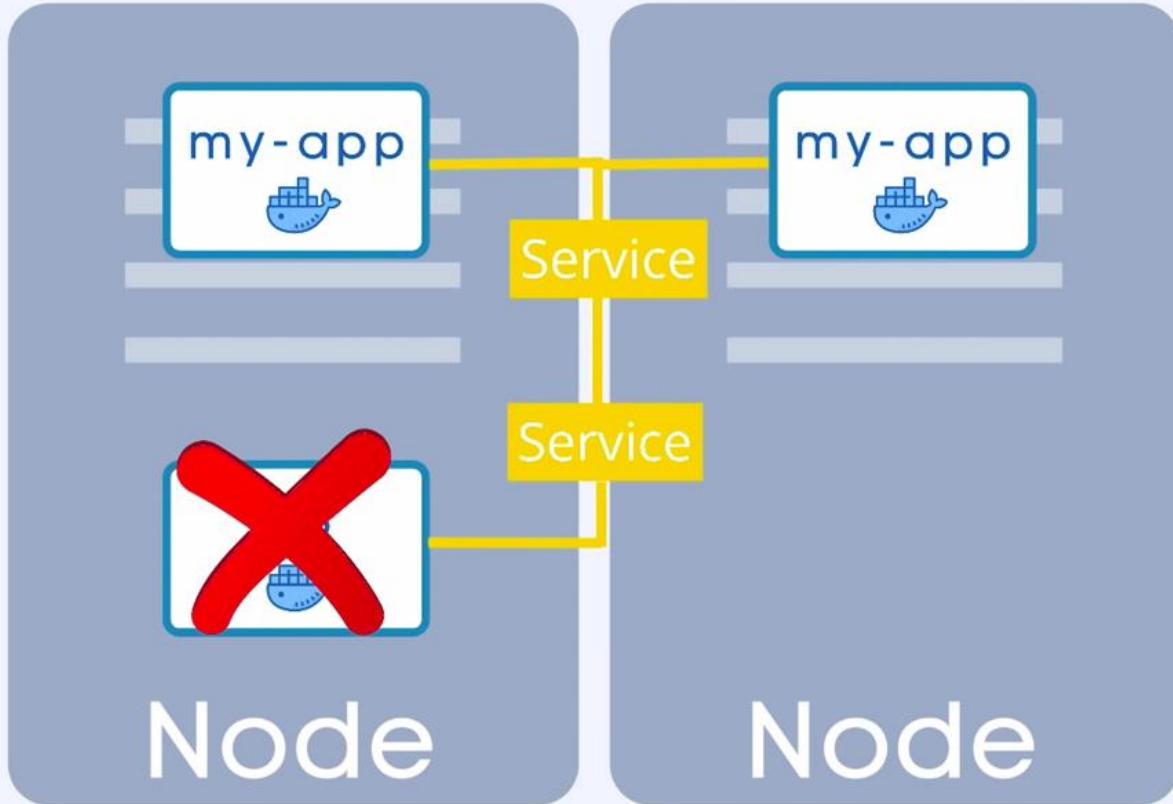
## DEPLOYMENT

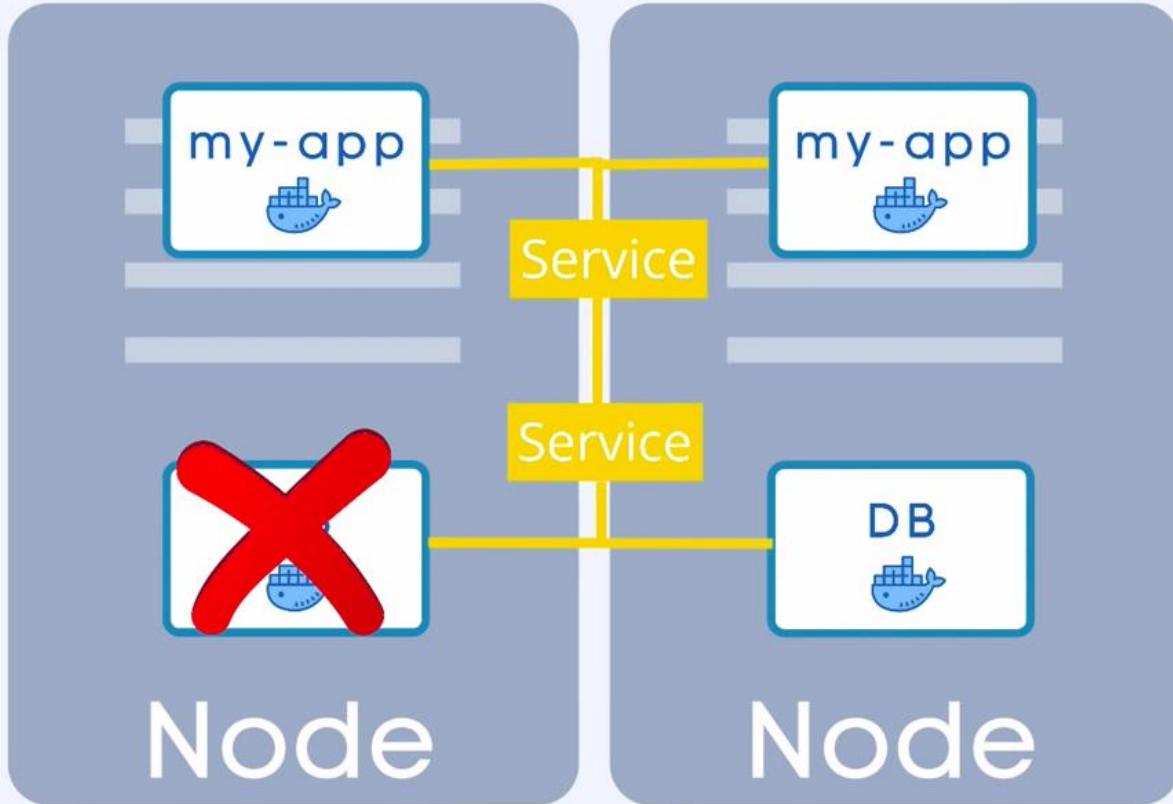


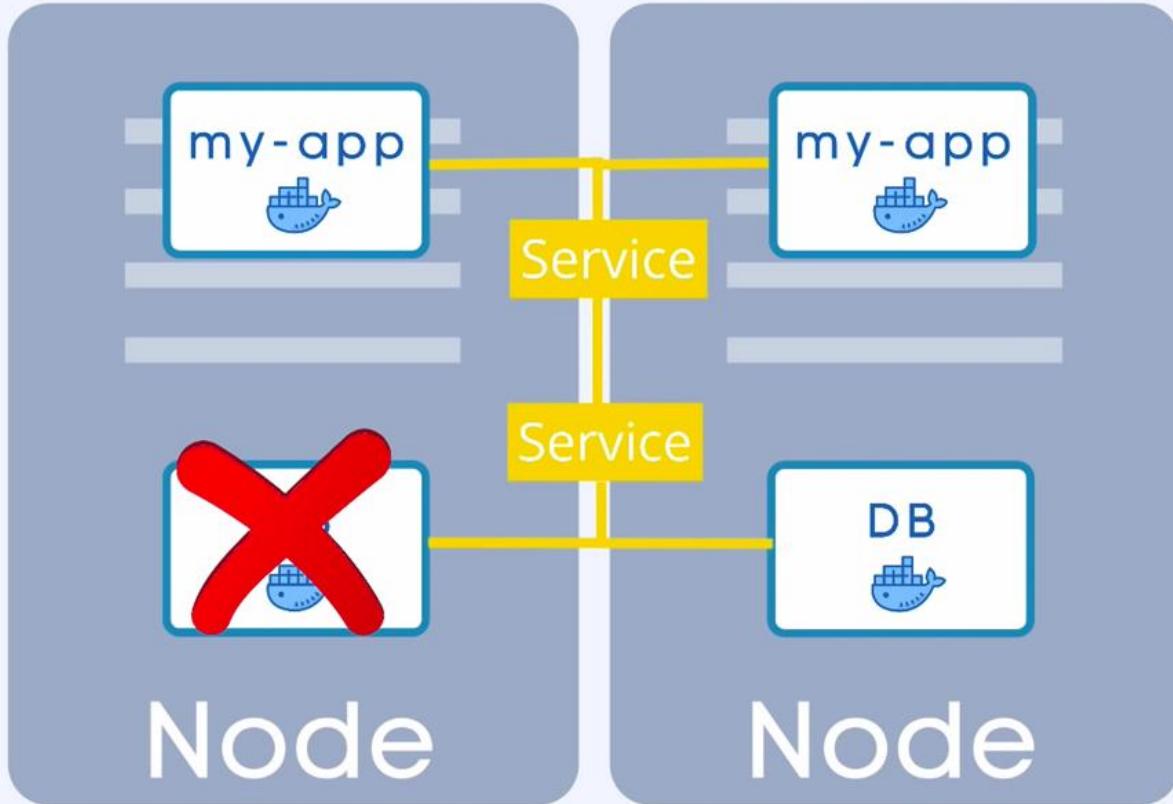
- ▶ Blueprint for "my-app" Pods
- ▶ You create Deployments
- ▶ Abstraction of Pods



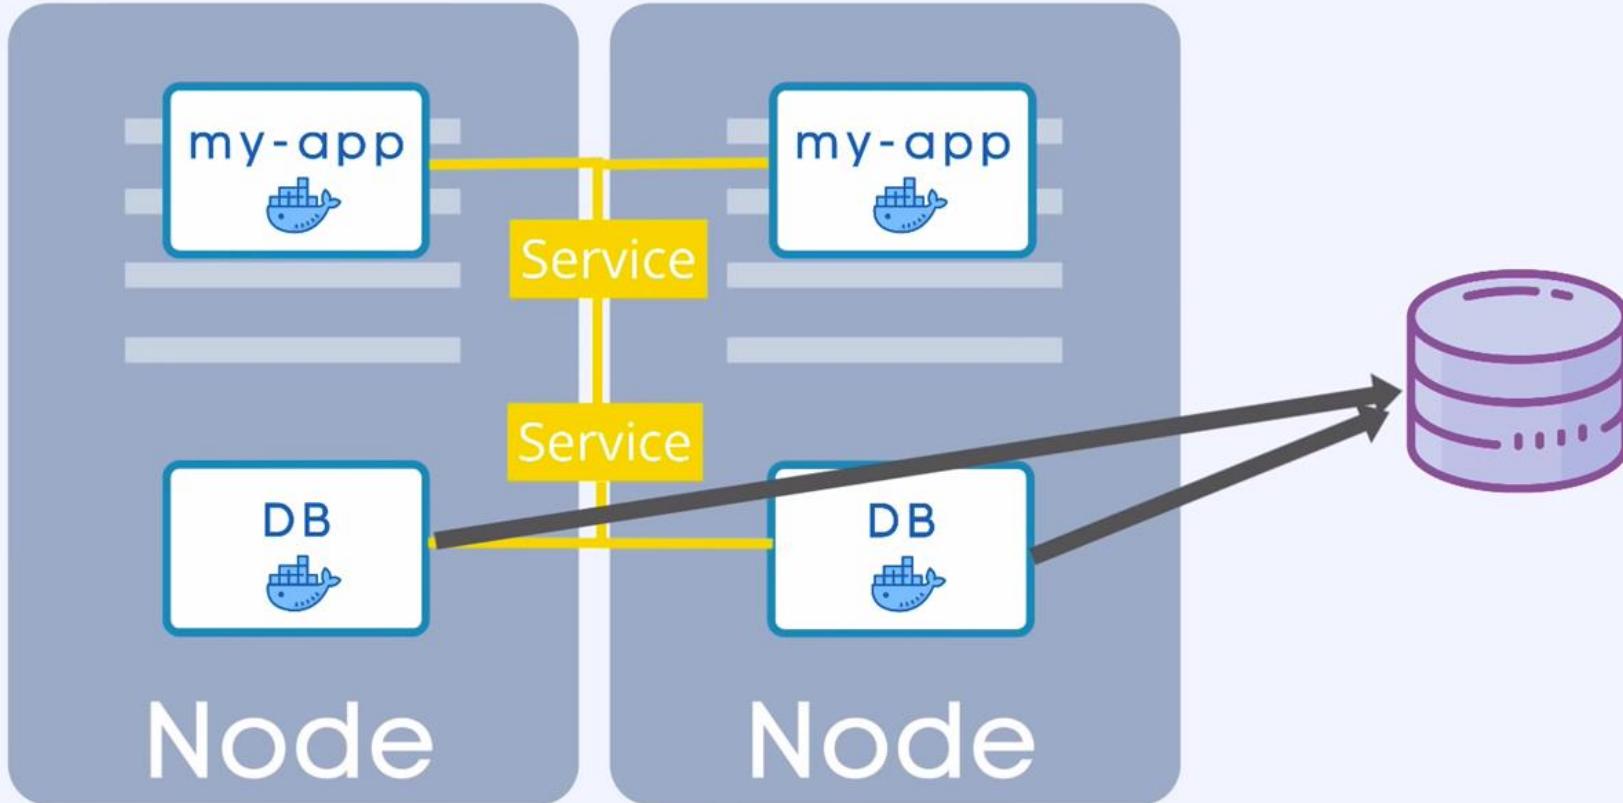


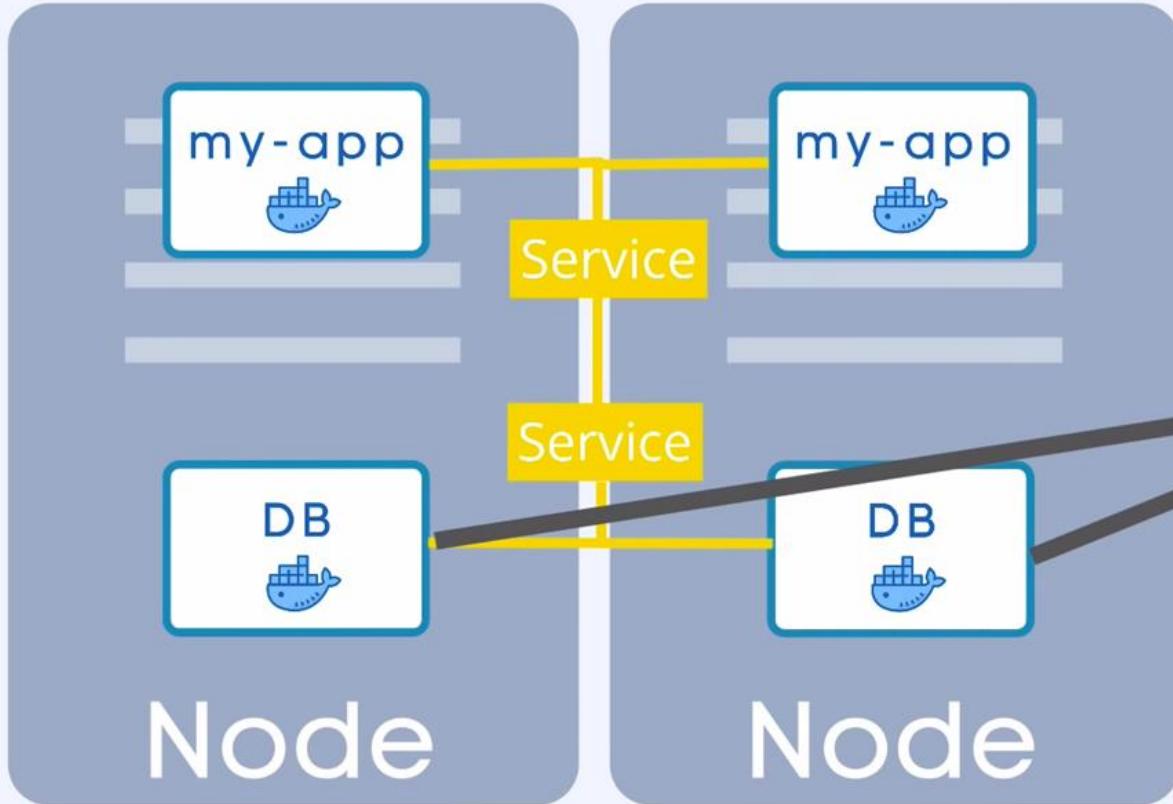




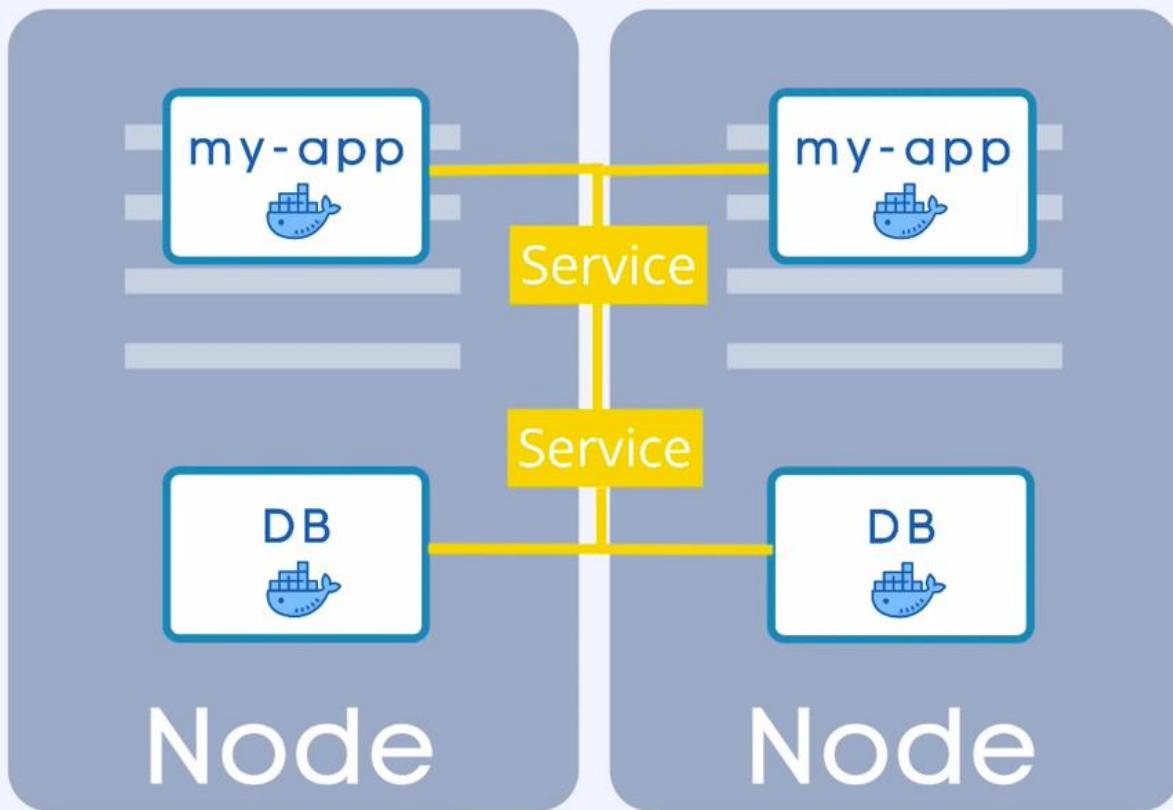


! DB can't be replicated via Deployment!



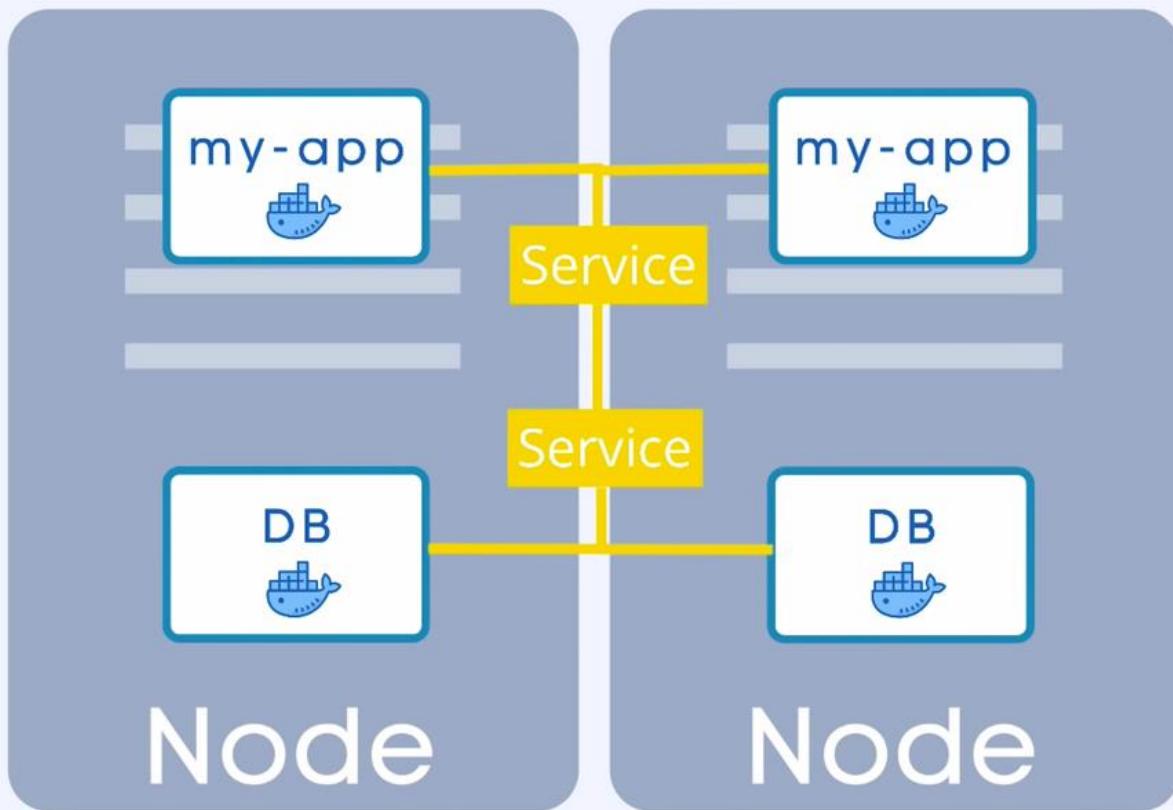


! Avoid Data Inconsistencies



STATEFULSET

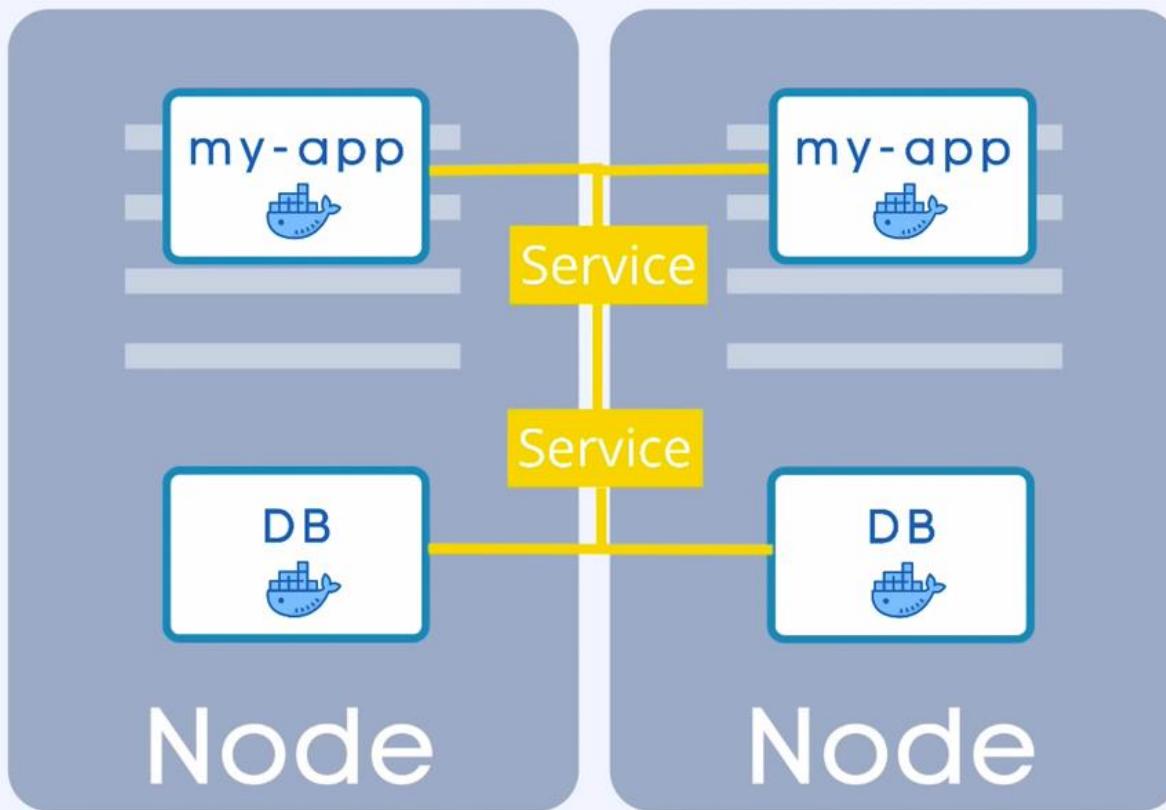




## STATEFULSET



► for  
STATEFUL apps

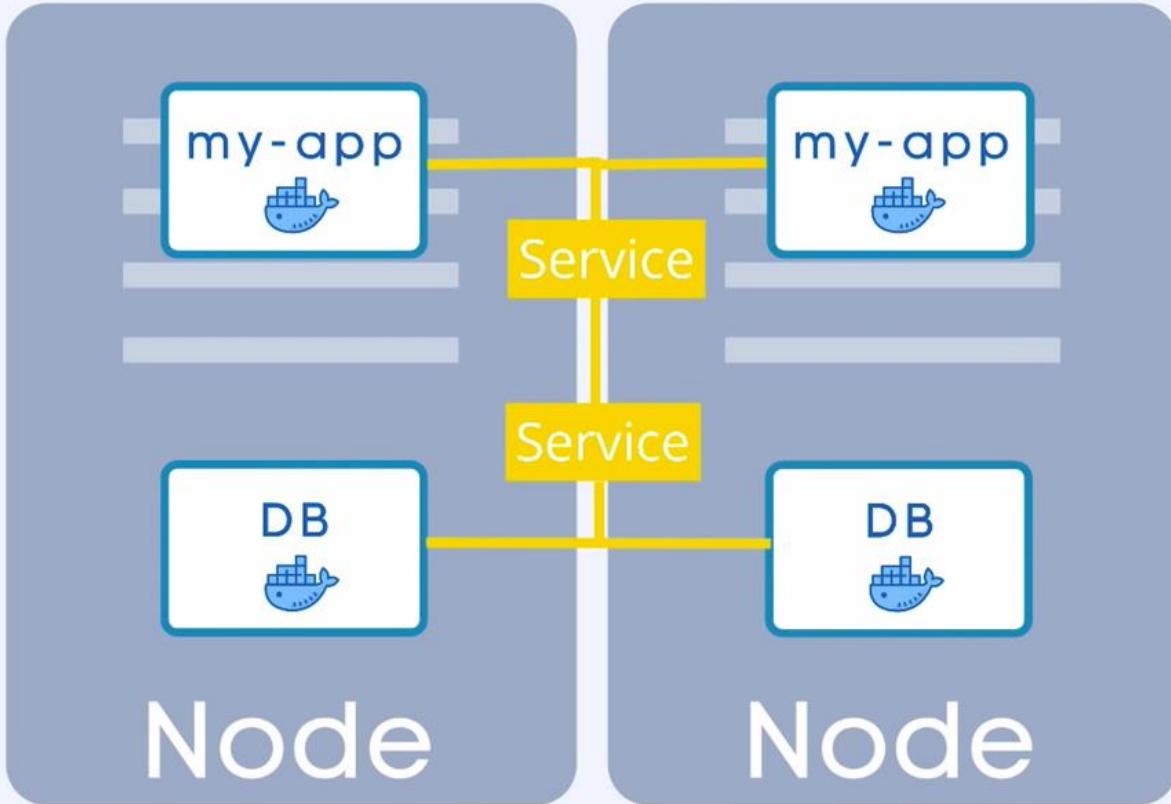


## STATEFULSET



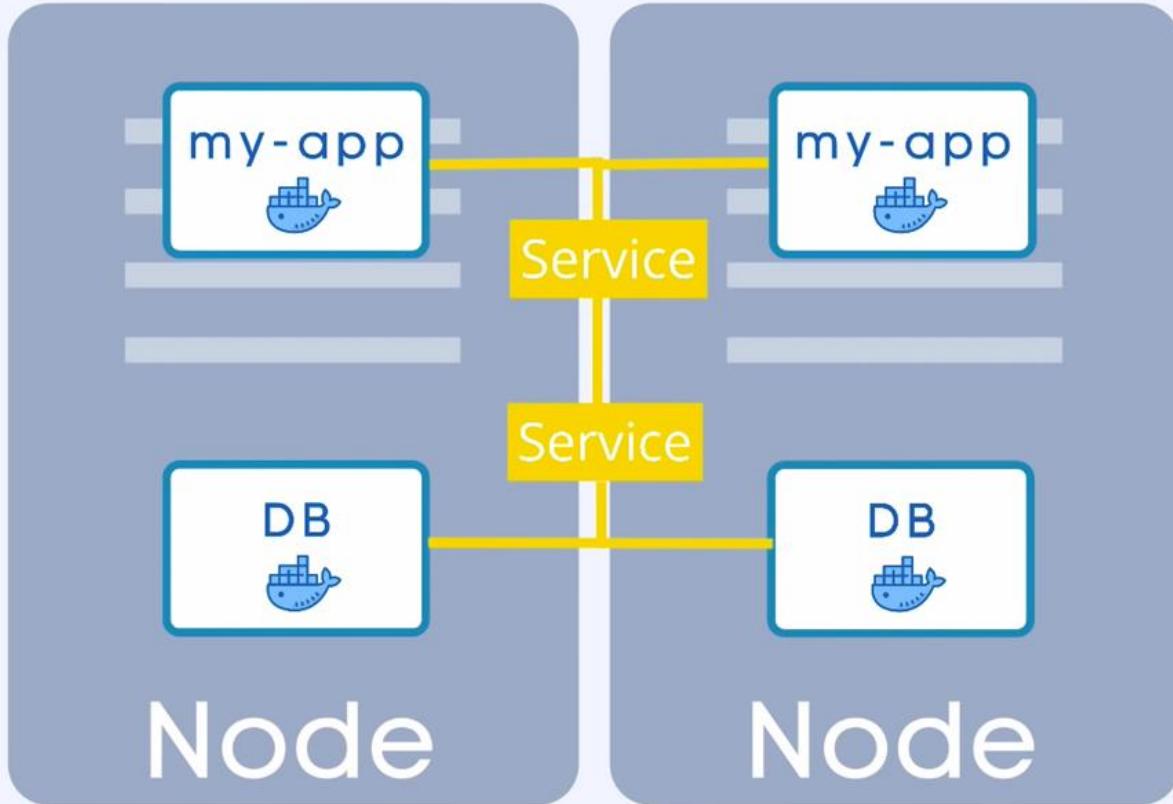
► for  
STATEFUL apps





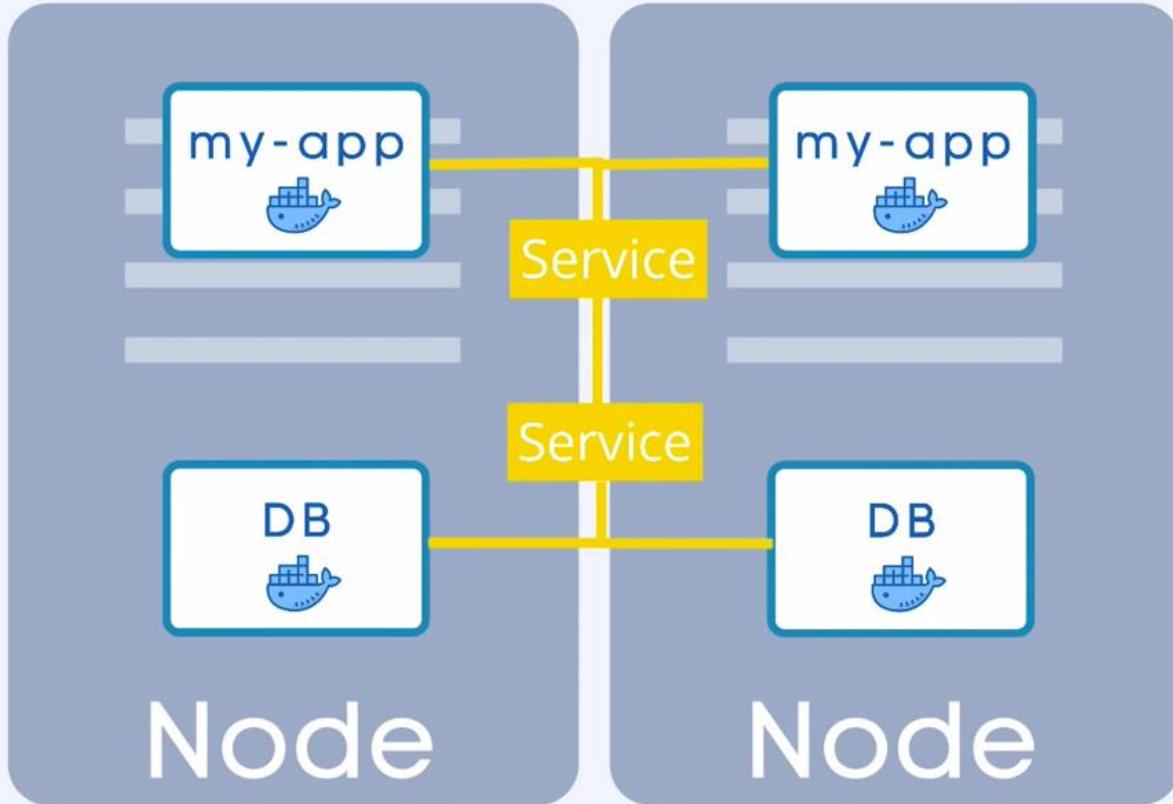
**Deployment** =  
for stateLESS Apps

**StatefulSet** =  
for stateFUL Apps or Databases

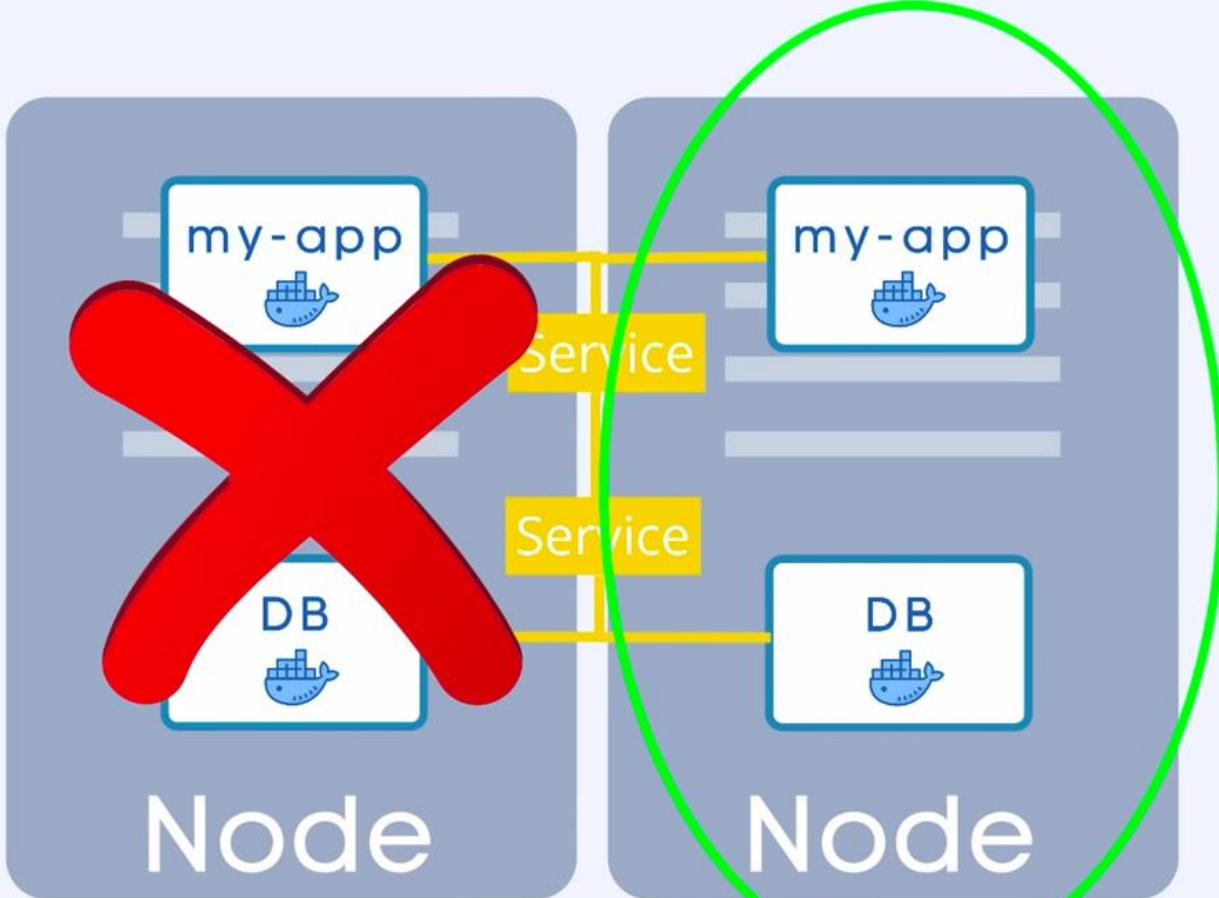


Deploying StatefulSet  
not easy

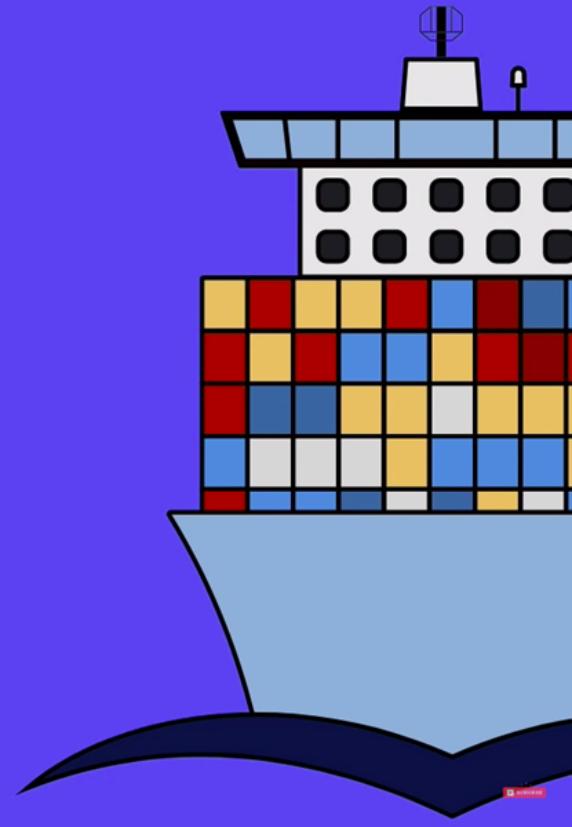




DB are often hosted outside  
of Kubernetes cluster



# Wrap Up





# Main Kubernetes Components

## summarized



Pod



abstraction of containers



Service



communication



Ingress



route traffic into cluster



# Main Kubernetes Components summarized



ConfigMap



Secret



external configuration



# Main Kubernetes Components

## summarized



Volume



data persistence



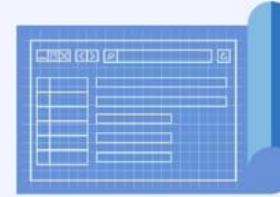
# Main Kubernetes Components summarized



Deployment



StatefulSet



★ replication